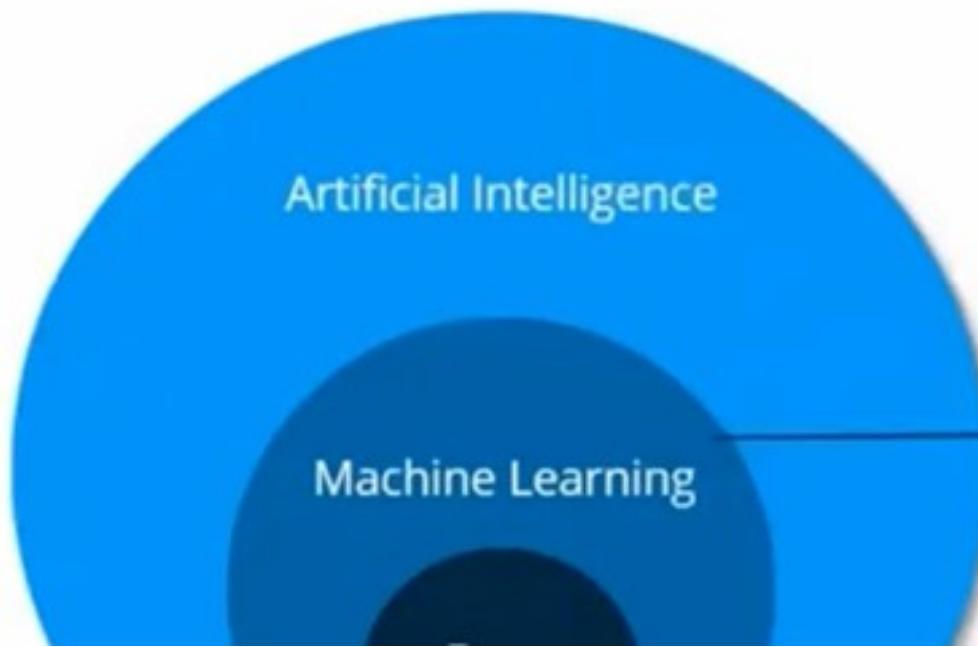


Deep Learning

taroonreddy.com

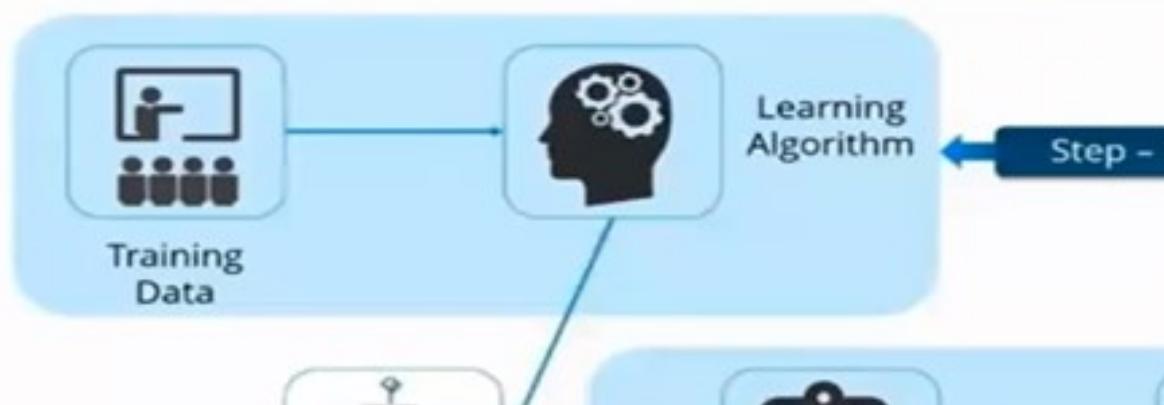


We'll learn more about Deep learning networks and Neural Networks

Machine Learning is a subset of AI

Types Of Machine Learning – Supervised

Supervised Learning is where you have input variables (x) and an output variable (y). The algorithm learns from the training data to predict the output for new data.



Types Of Machine Learning – Unsupervised

- ❑ Unsupervised Learning is the training of a model using information that is neither classified nor labelled.
- ❑ This model can be used to cluster the input data in classes on the basis of their statistical properties.



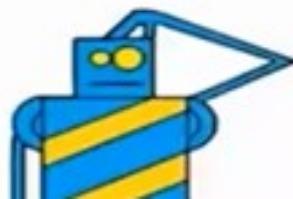
Types Of Machine Learning – Reinforcement Learning

- ❑ Reinforcement Learning (RL) is learning by interacting with a simulated environment.
- ❑ An RL agent learns from the consequences of its actions, rather than being explicitly told what to do. It selects its actions on basis of its past experiences (trial and error), and explores different choices (exploration).



Limitations Of Machine Learning

- ❑ One of the big challenges with traditional Machine Learning is feature extraction.
- ❑ For complex problems such as object recognition or handwriting, challenge.



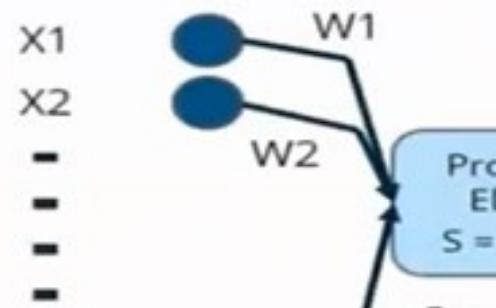
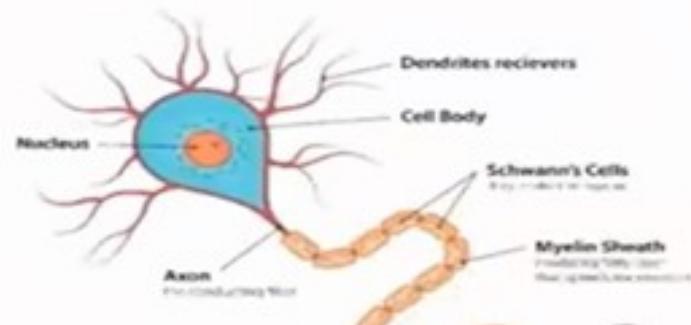
Deep Learning To The Rescue

- Deep Learning models are capable to focus on the right features with little guidance from the programmer.
- These models also partially solve the dimensionality problem.



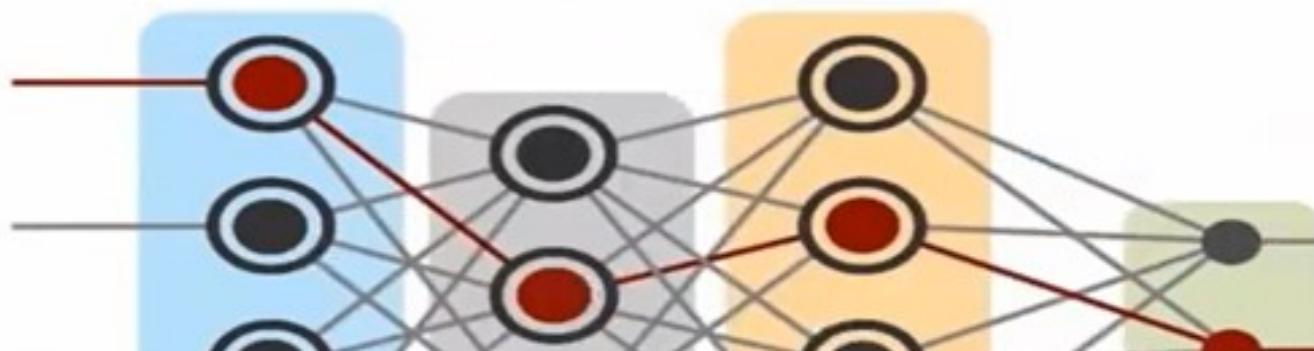
Deep Learning To The Rescue

- Deep Learning is implemented through Neural Networks.
- Motivation behind Neural Networks is the biological Neuron.



What Is Deep Learning?

A collection of statistical machine learning techniques used to perform tasks often based on artificial neural networks



taroonreddy.com

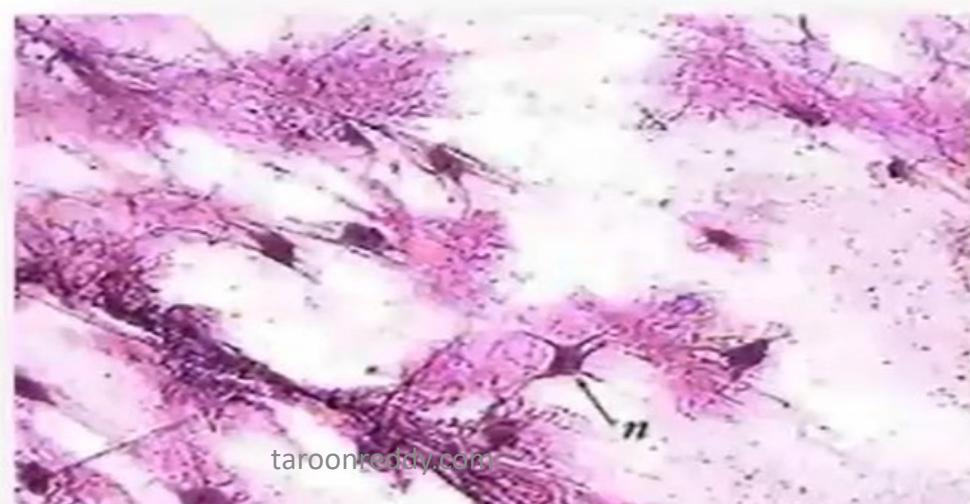
Neural Network

taroonreddy.com

The Neur

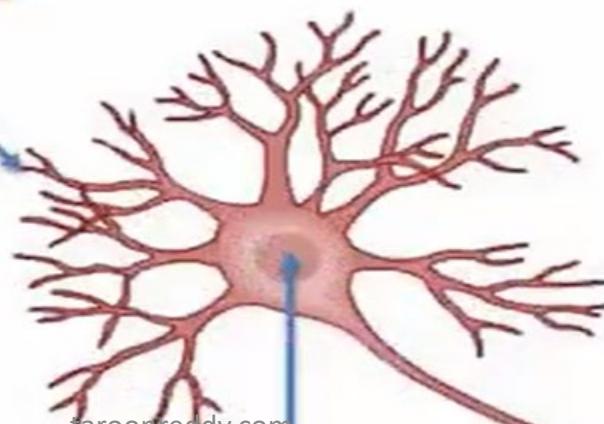
taroonreddy.com

The Neuron



The Neuron

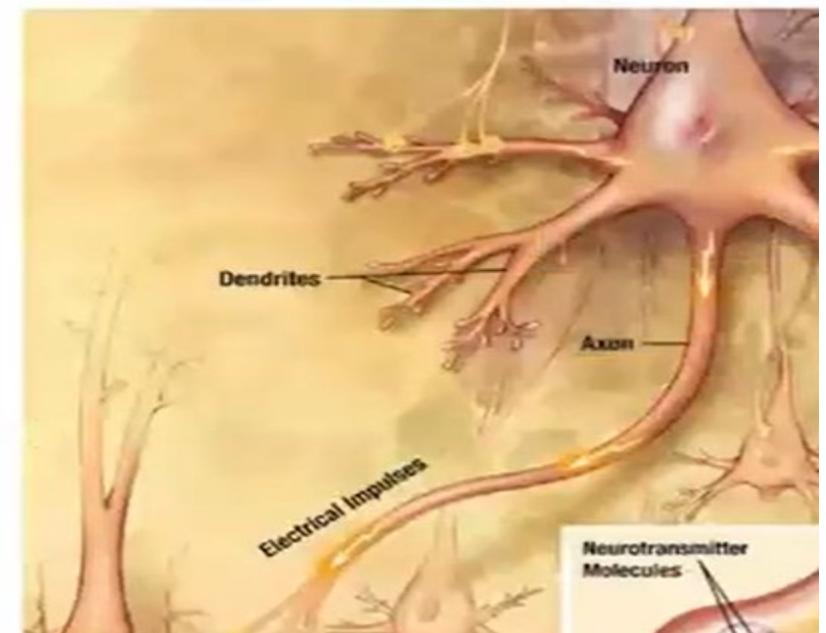
Dendrites



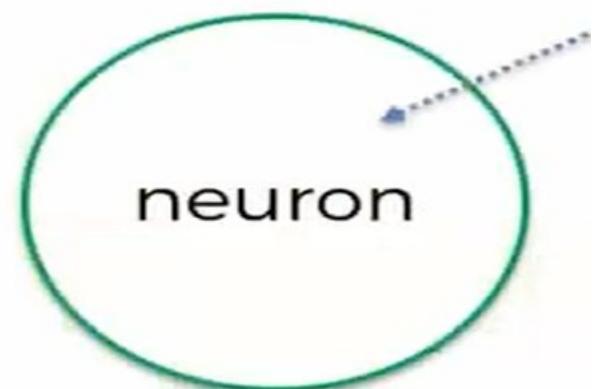
Ax

The Neuron

taroonreddy.com

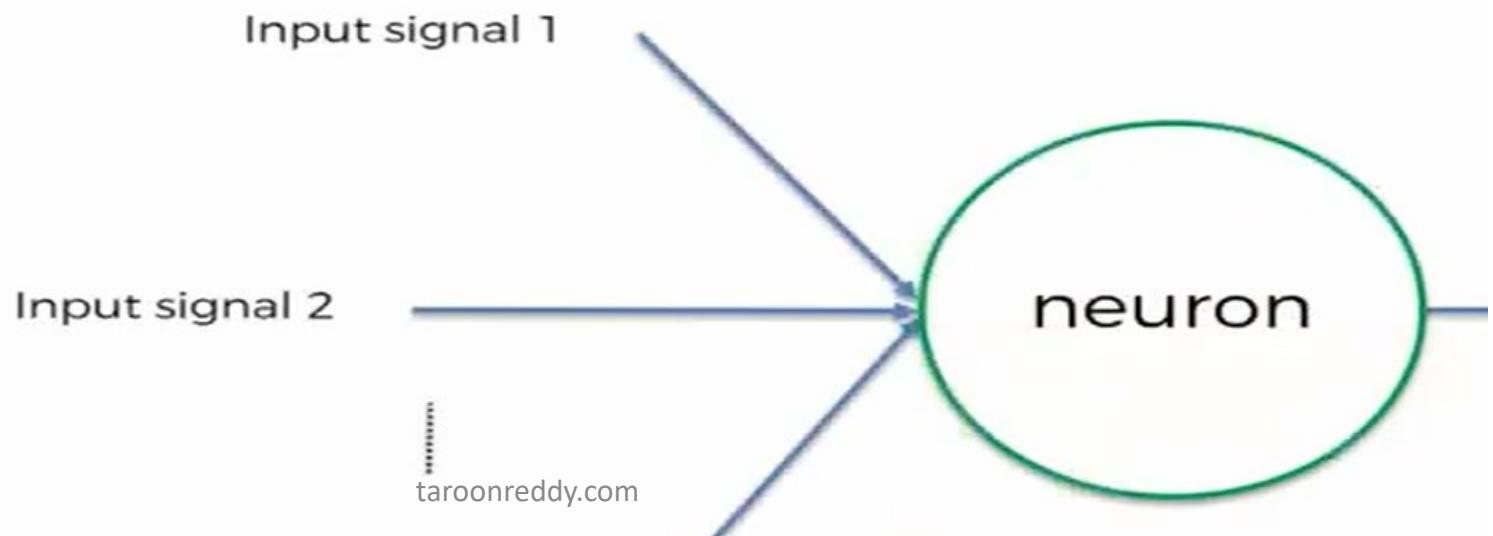


The Neuron

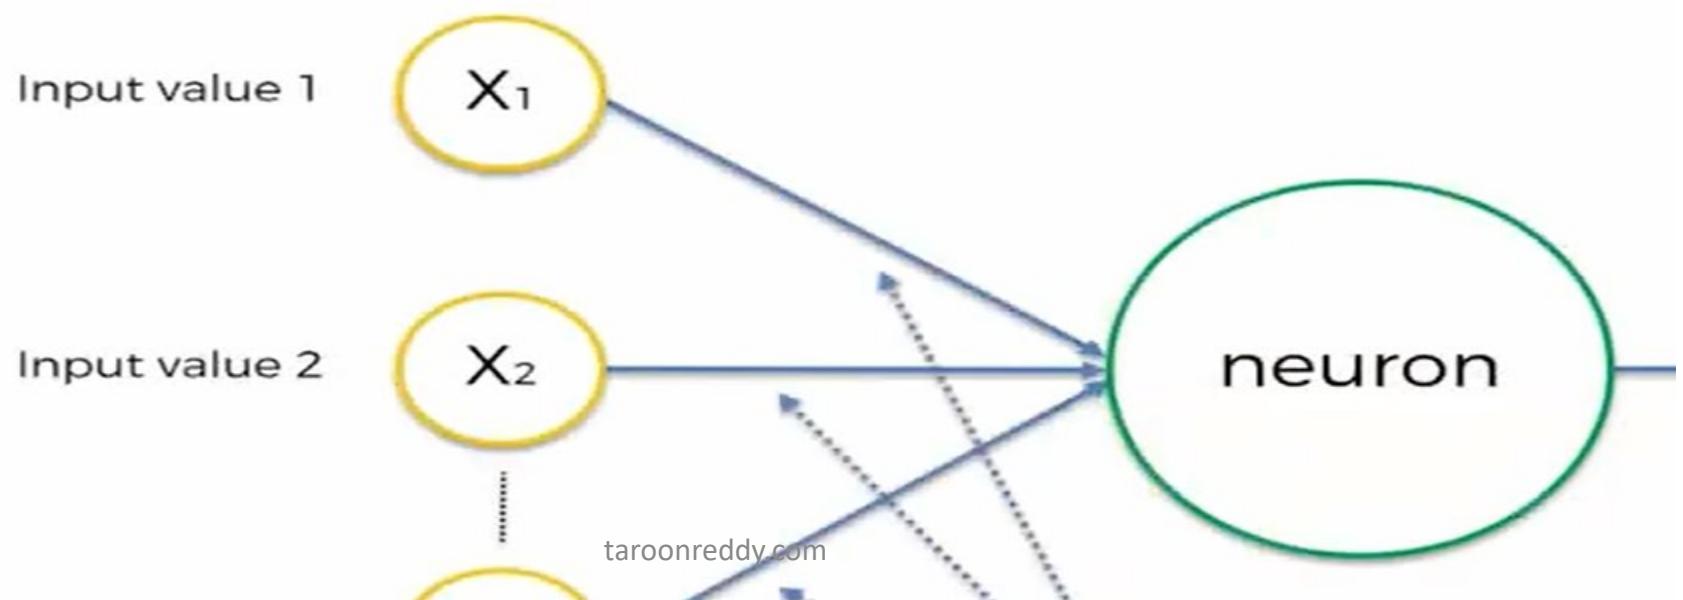


taroonreddy.com

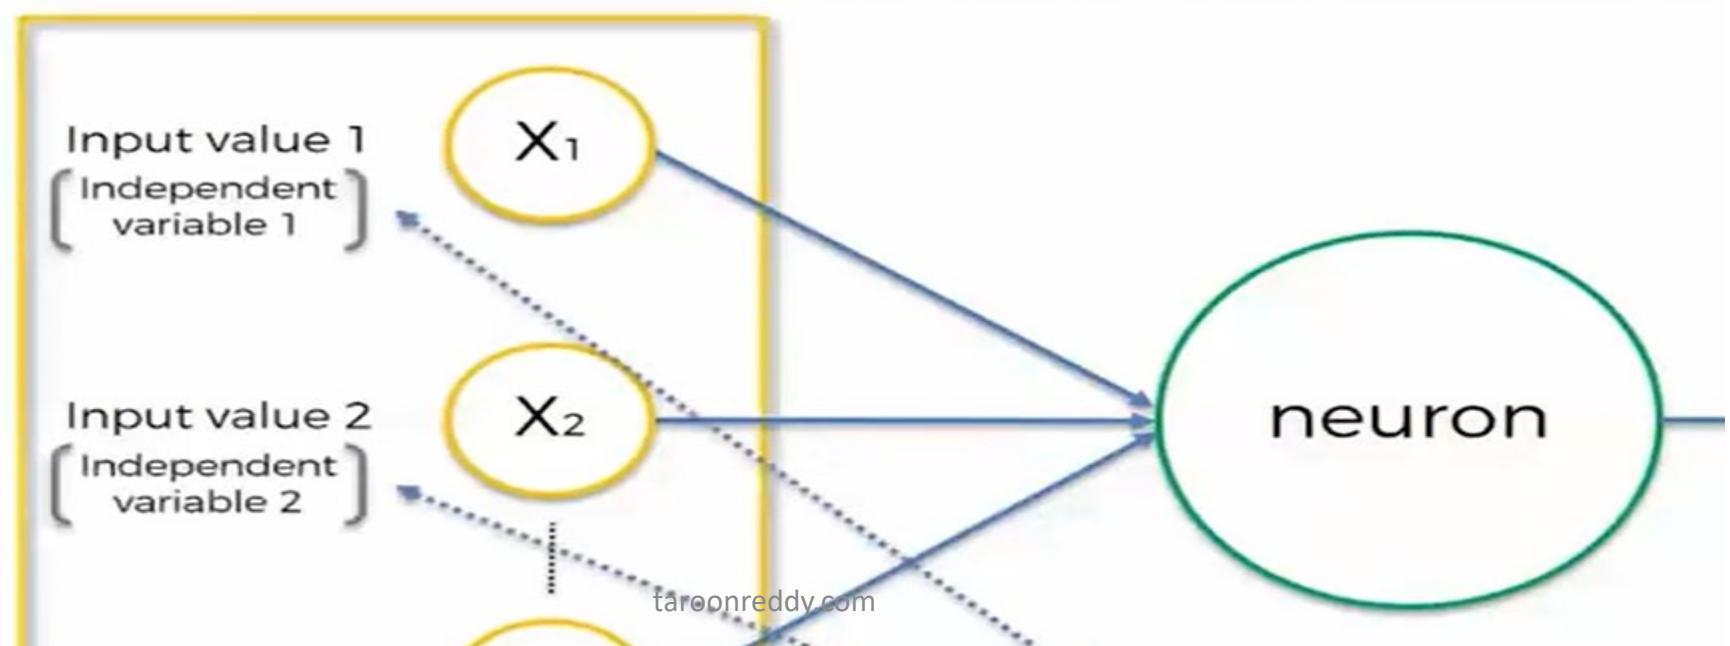
The Neuron



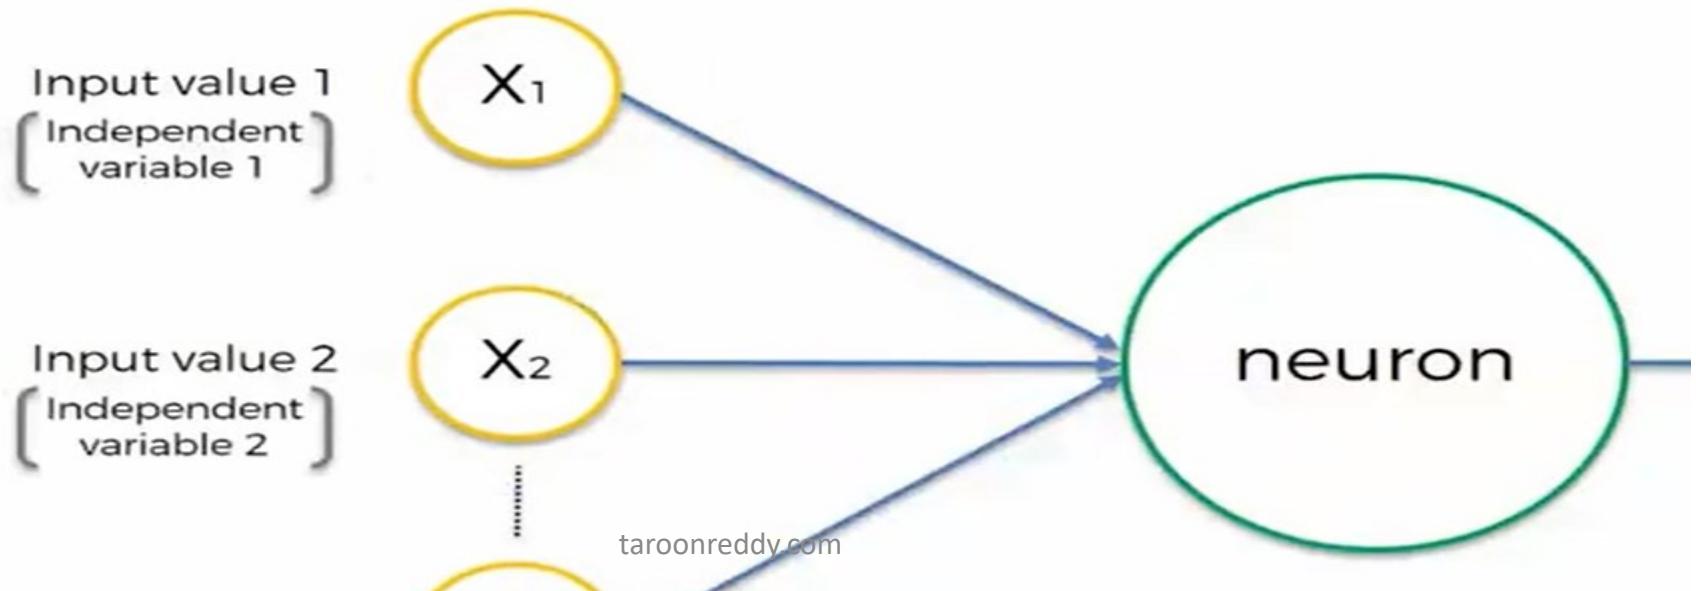
The Neuron



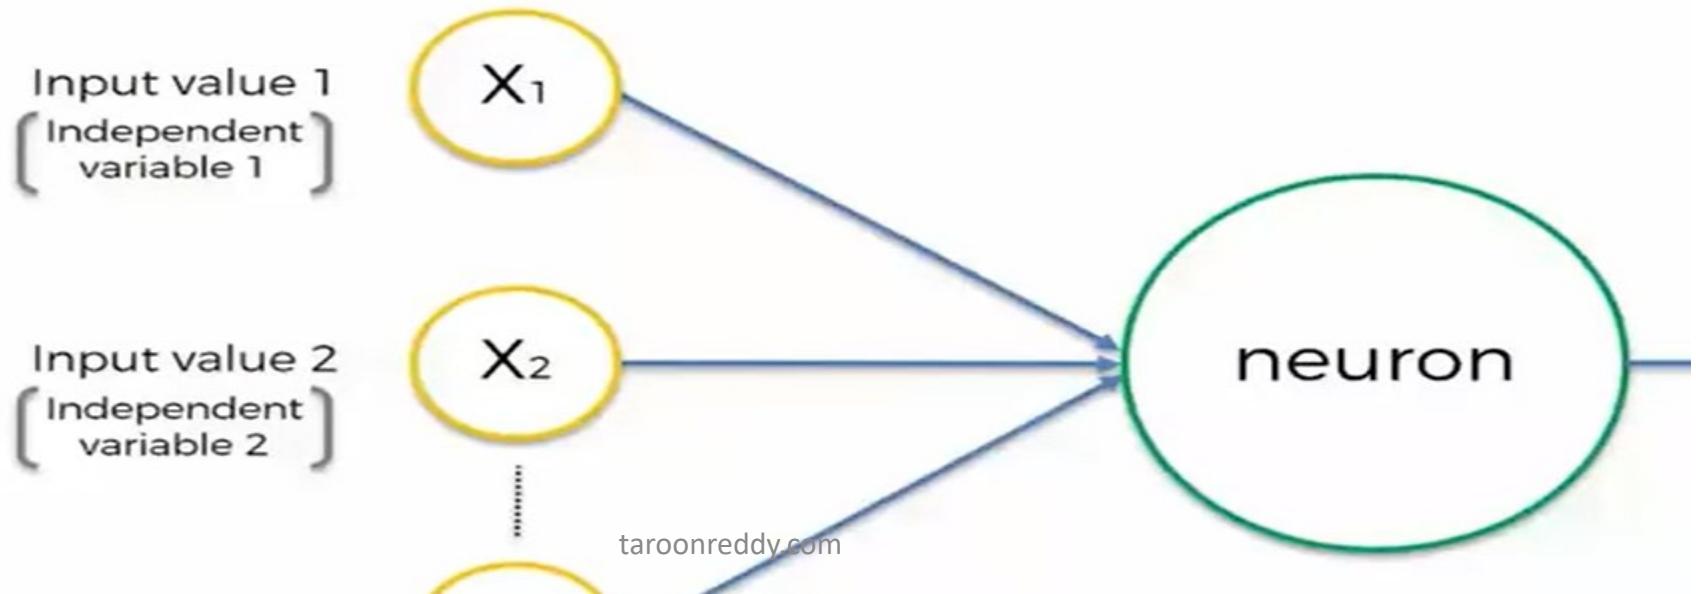
The Neuron



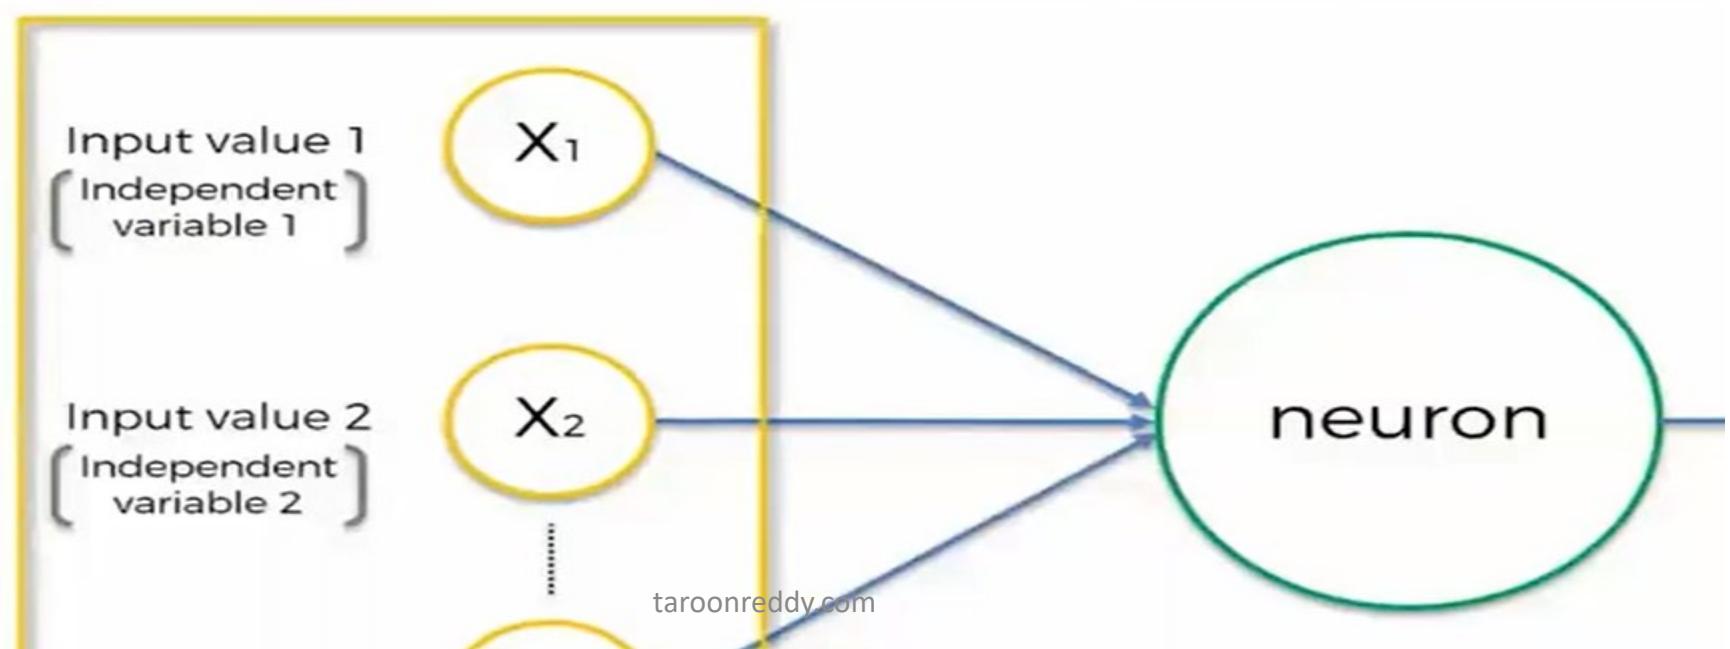
The Neuron



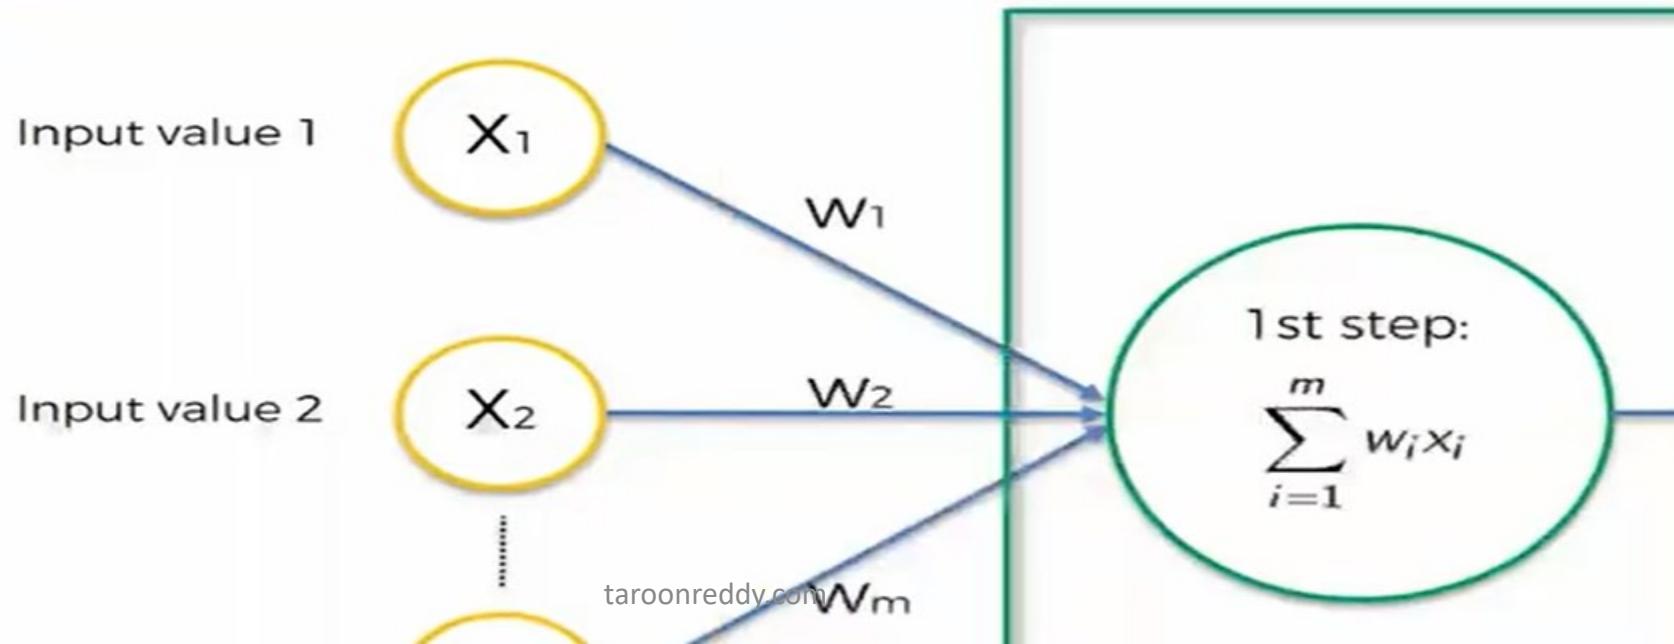
The Neuron



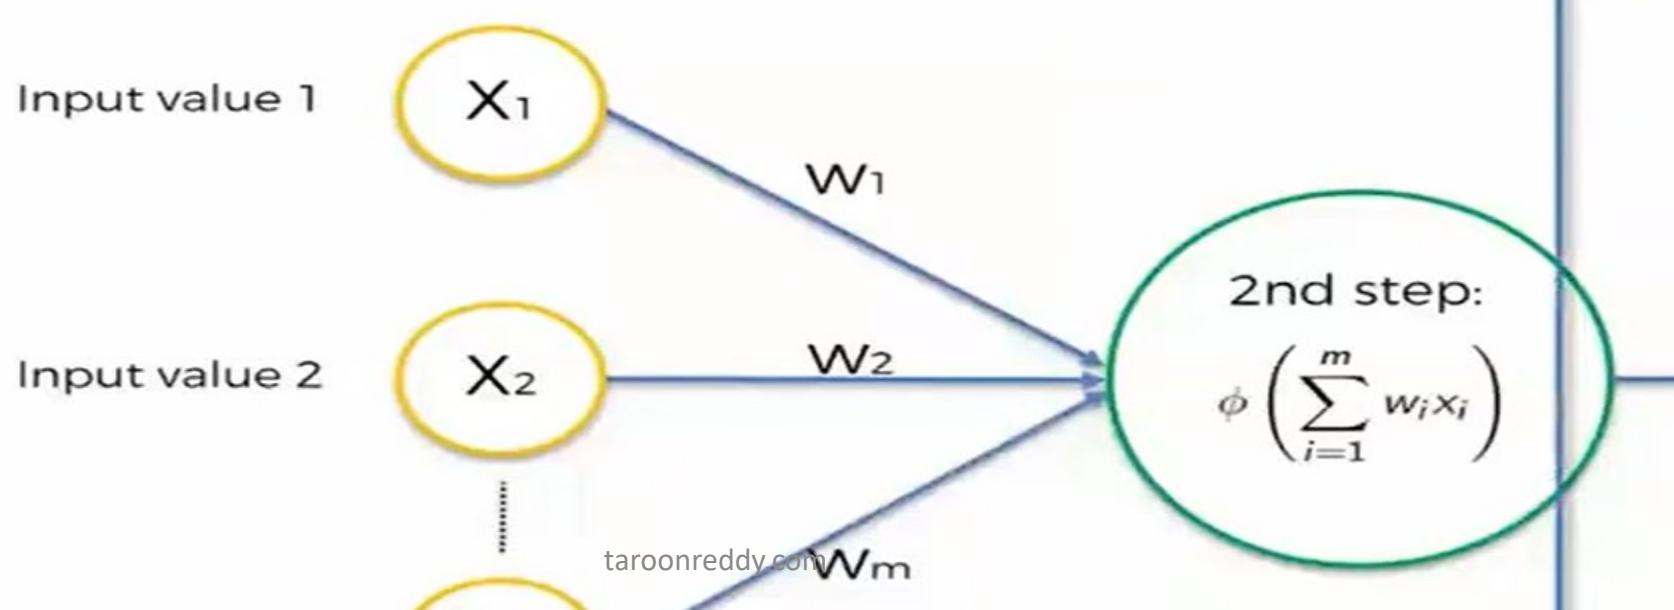
The Neuron



The Neuron



The Neuron



Activation Functions

The Activation Function

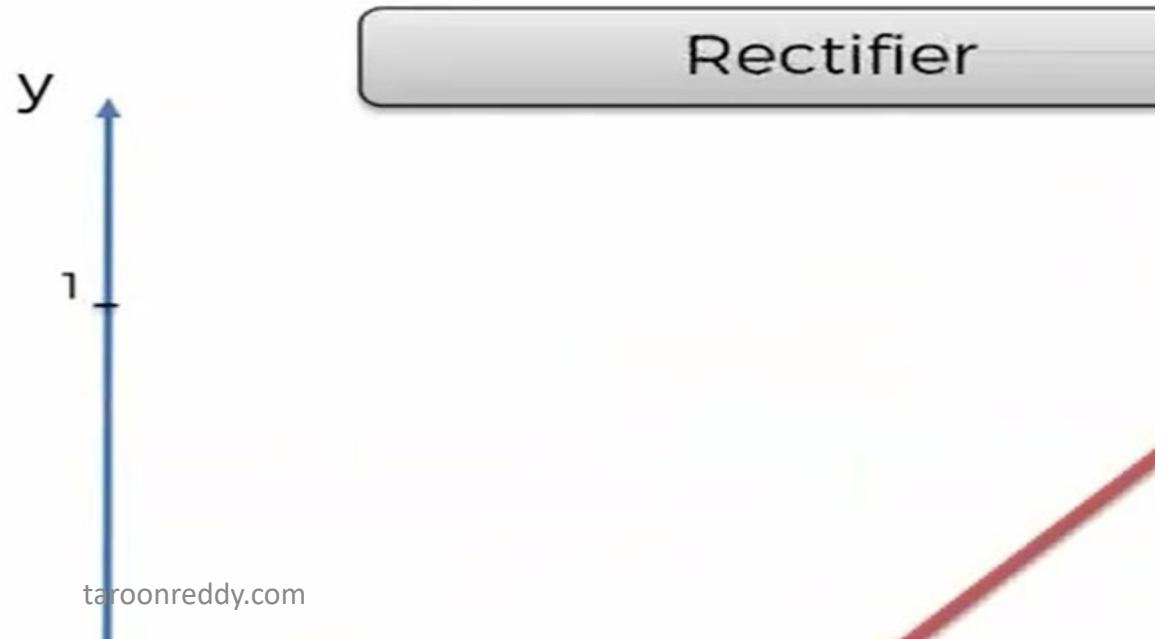
Threshold Function



The Activation Function

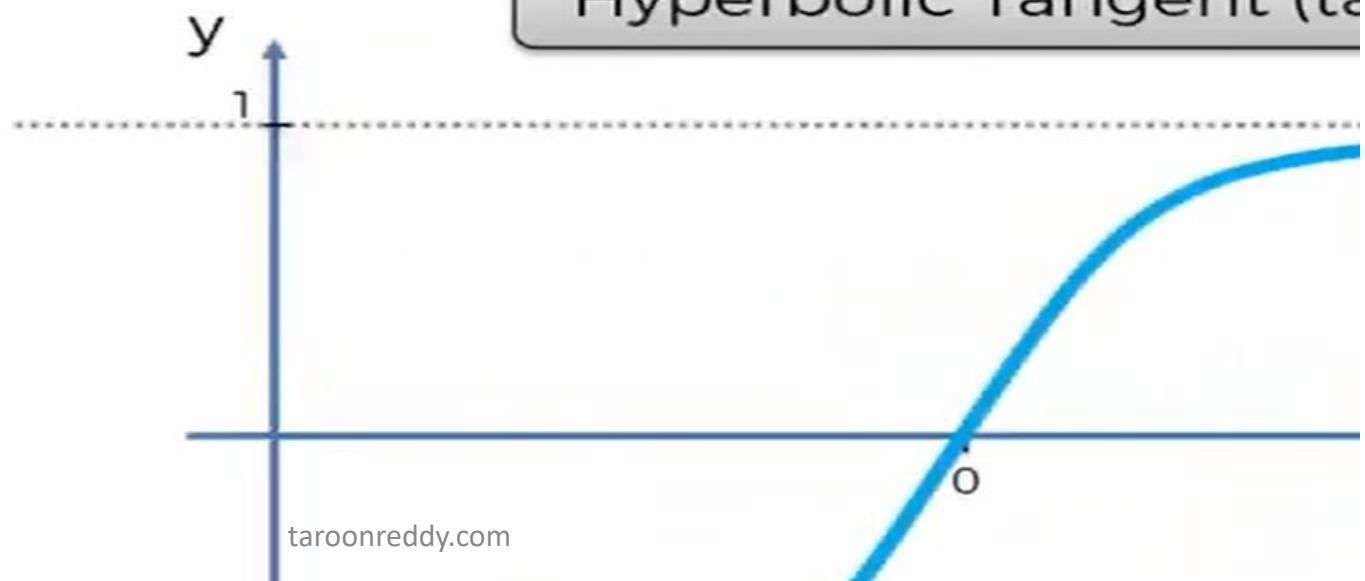


The Activation Function

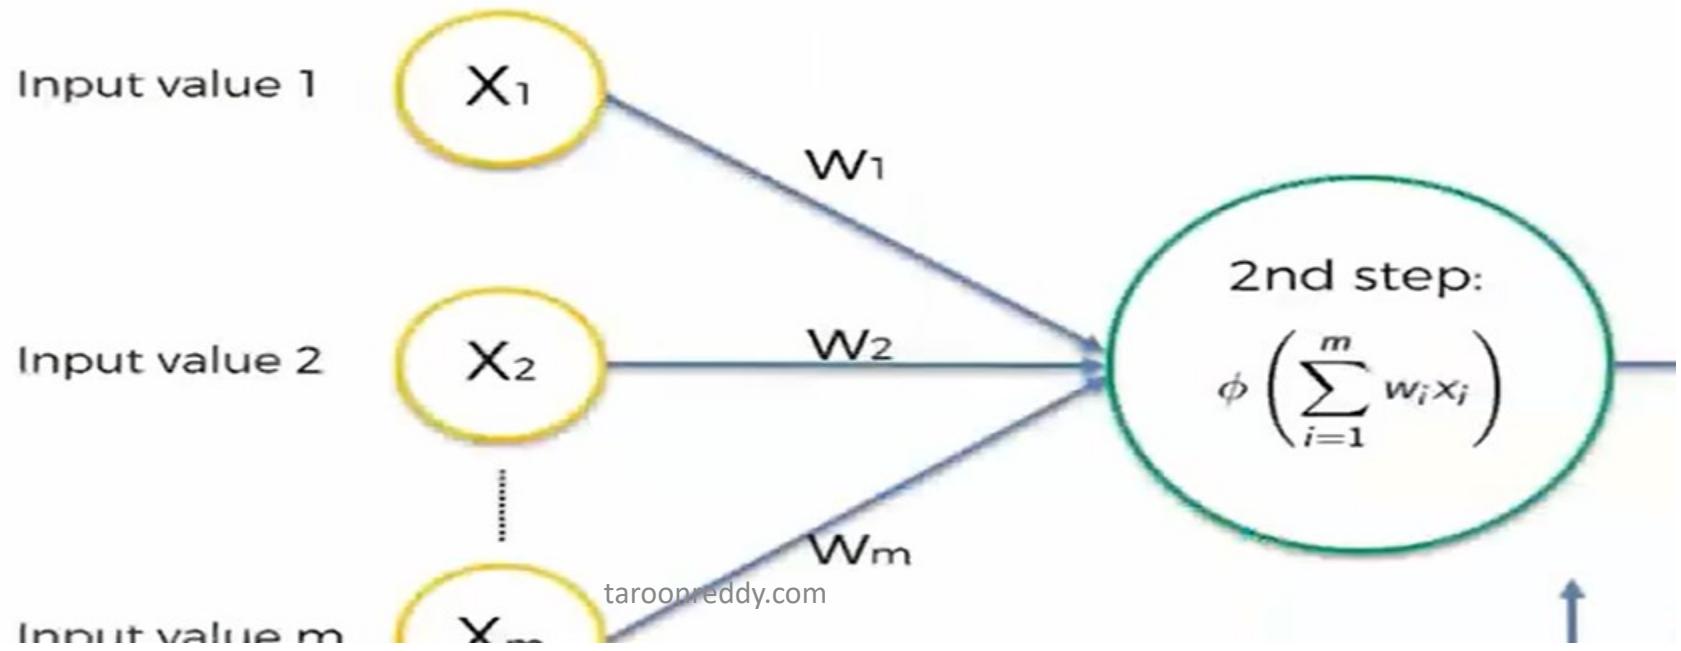


The Activation Function

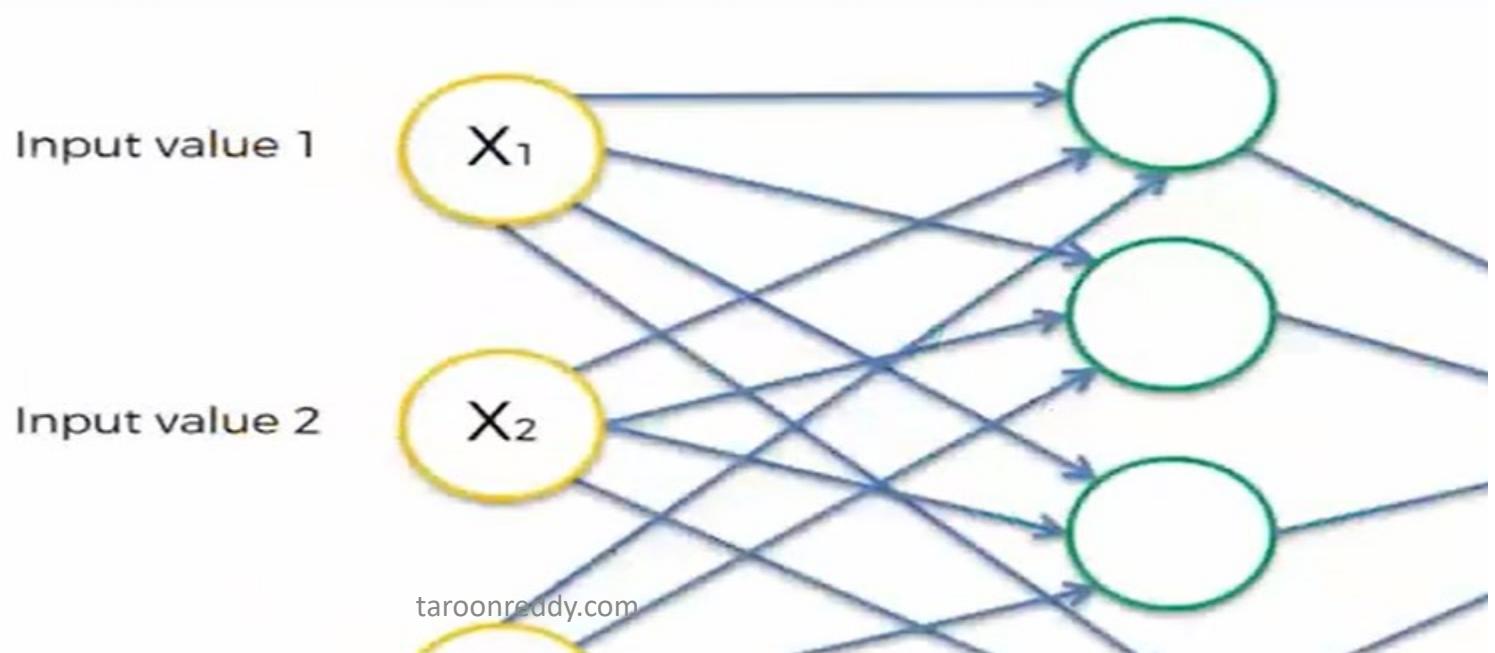
Hyperbolic Tangent (\tanh)



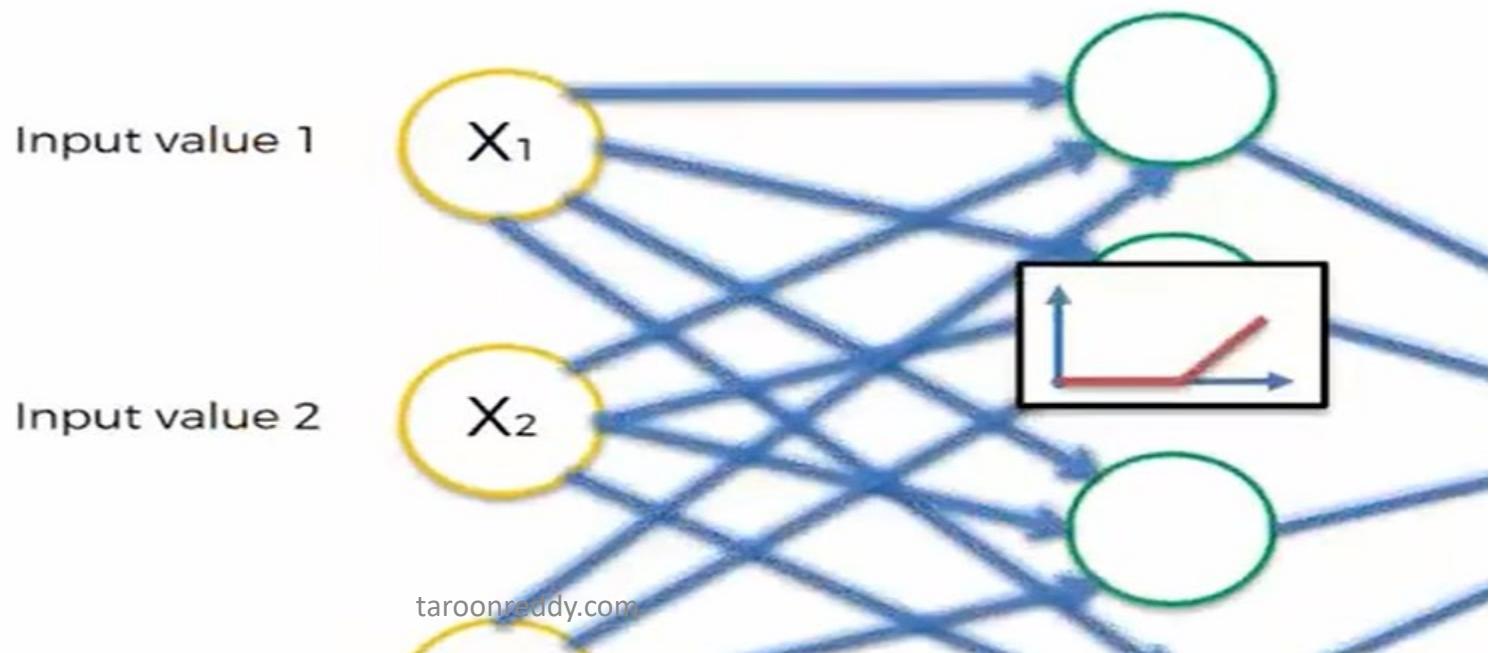
The Activation Function



The Activation Function



The Activation Function



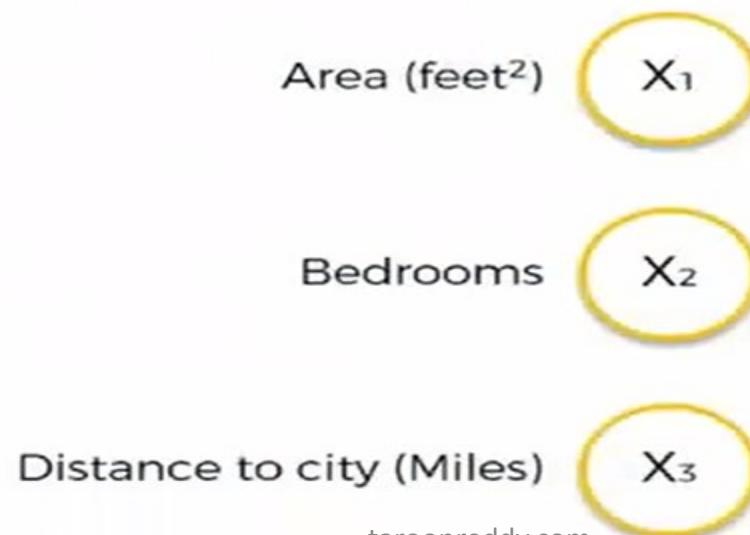
How do NNs

taroonreddy.com

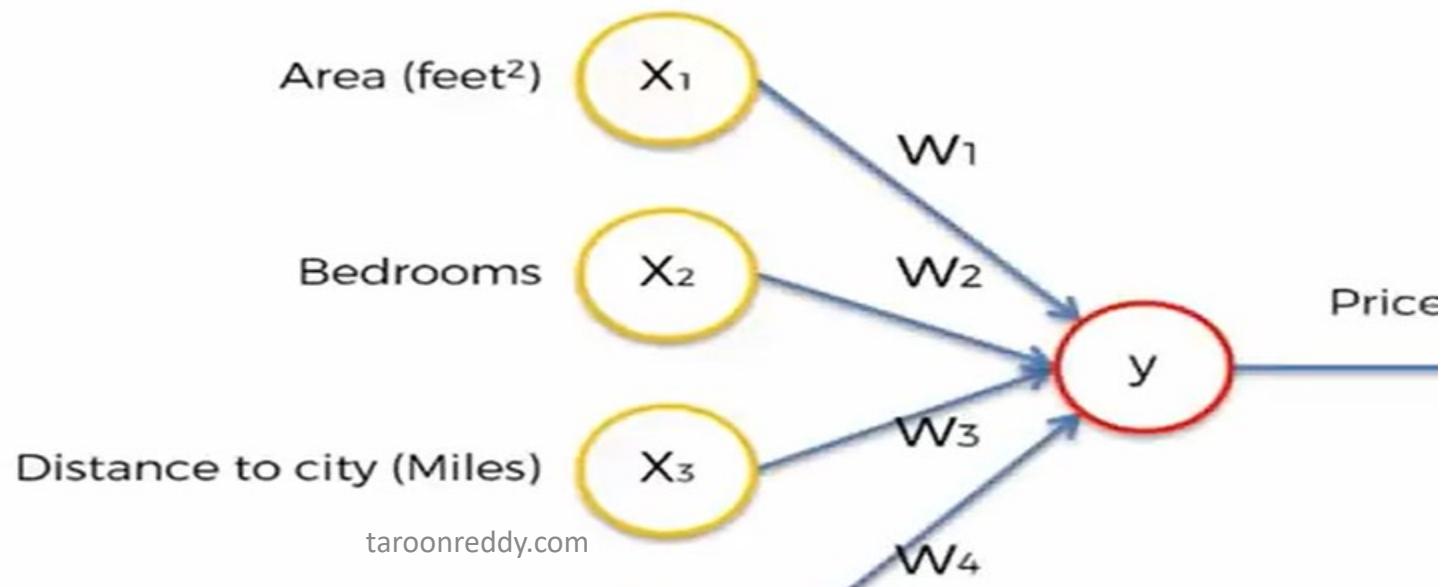


taroonreddy.com

How Do Neural Networks Work?



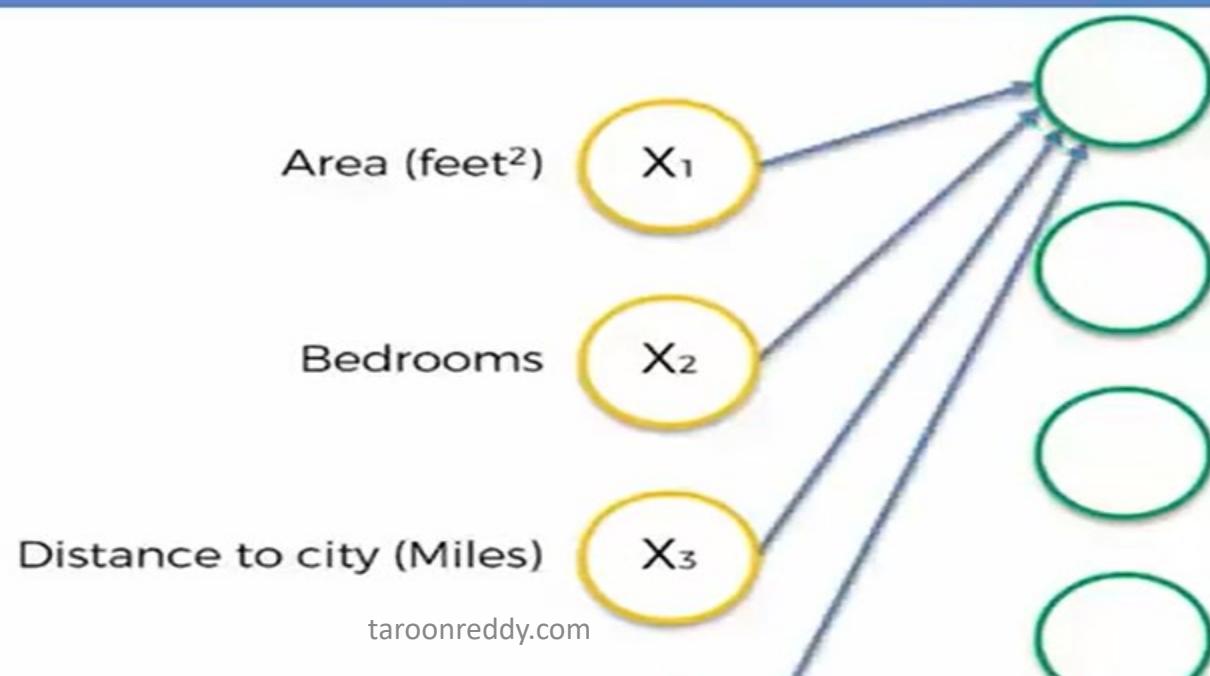
How Do Neural Networks Work?



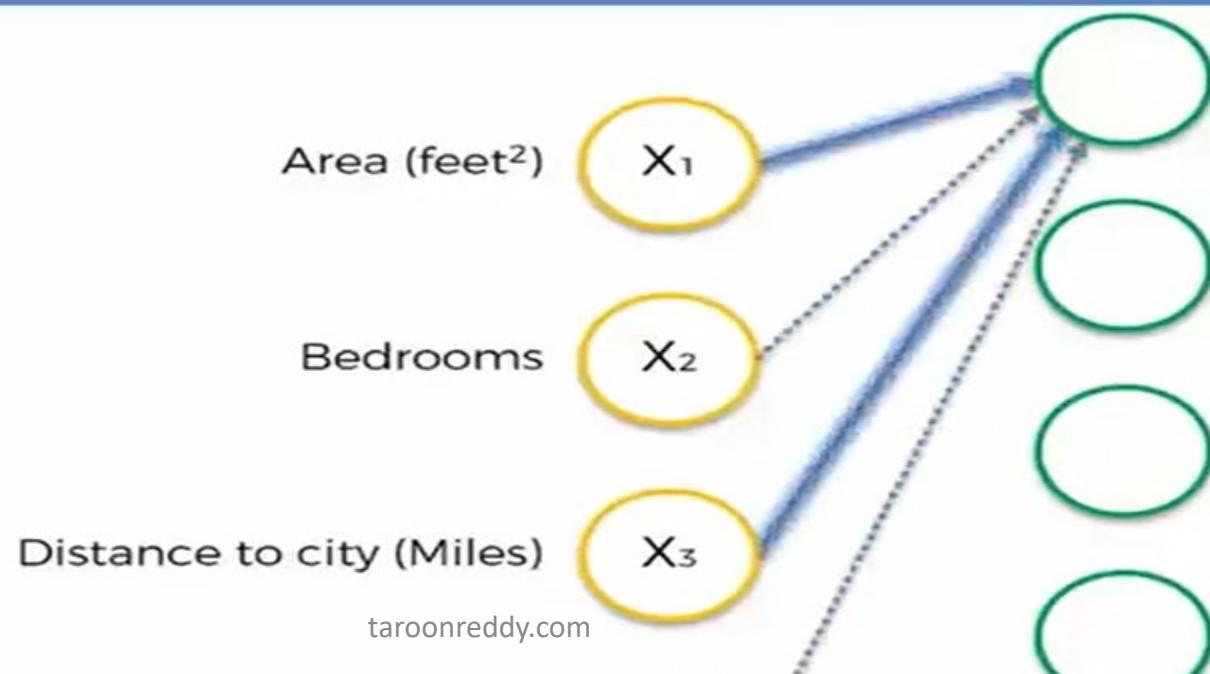
How Do Neural Networks Work?



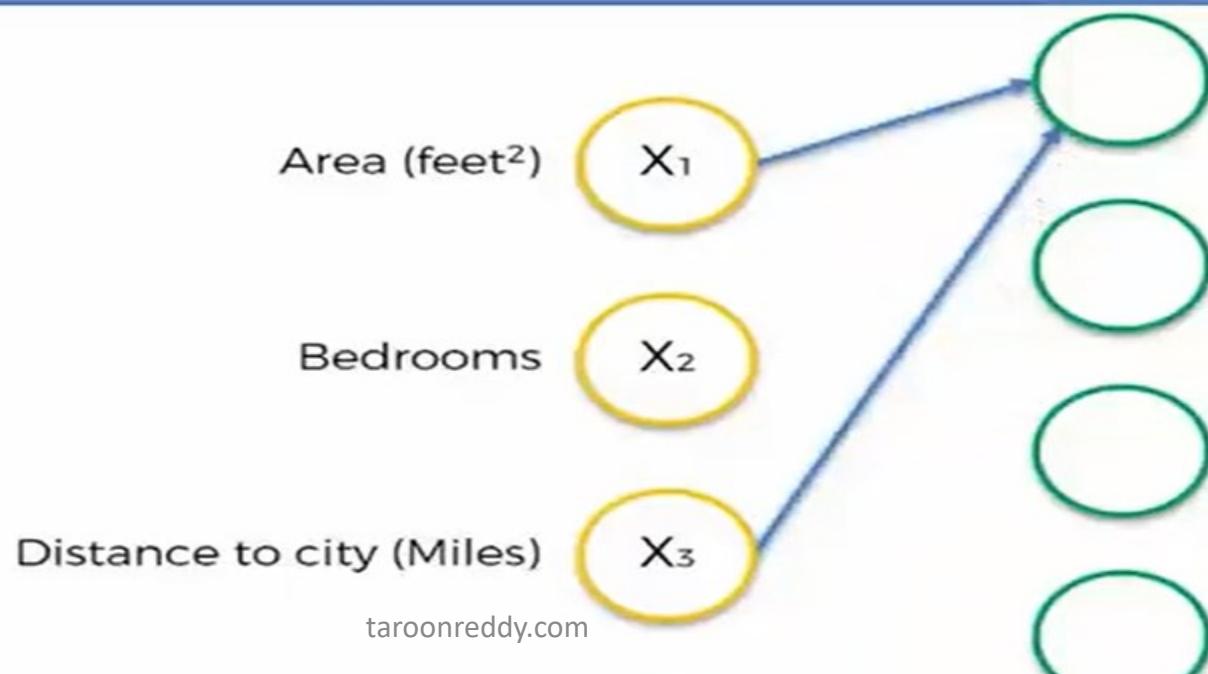
How Do Neural Networks Work?



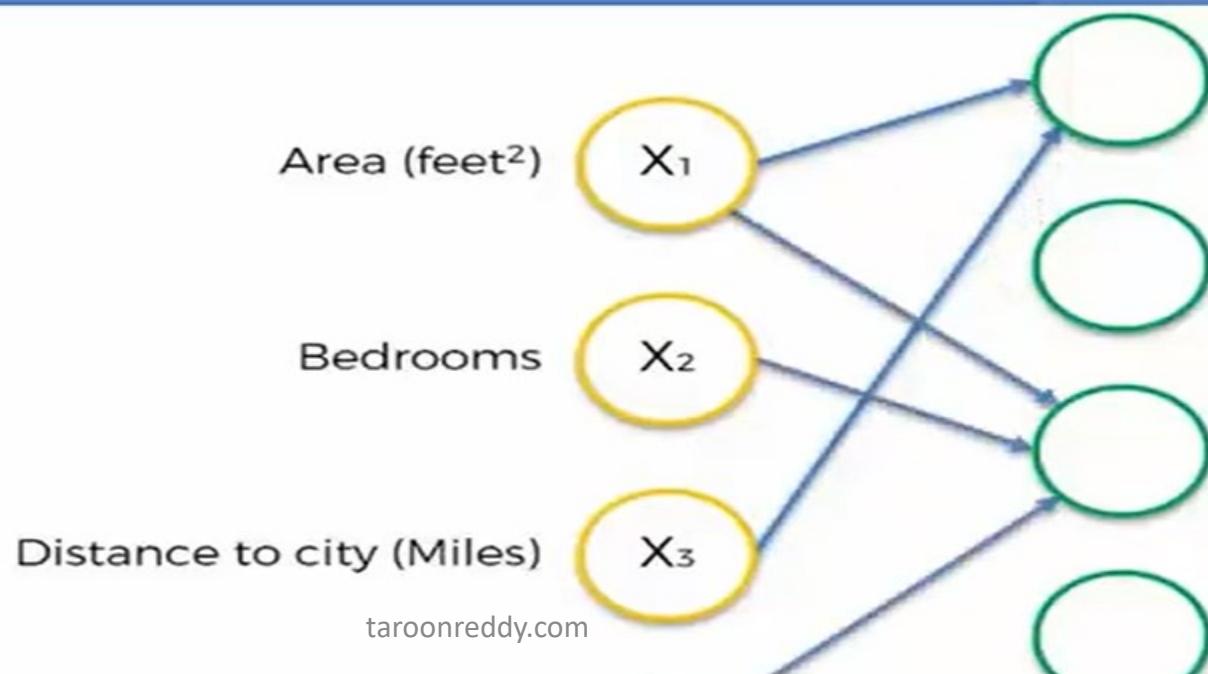
How Do Neural Networks Work?



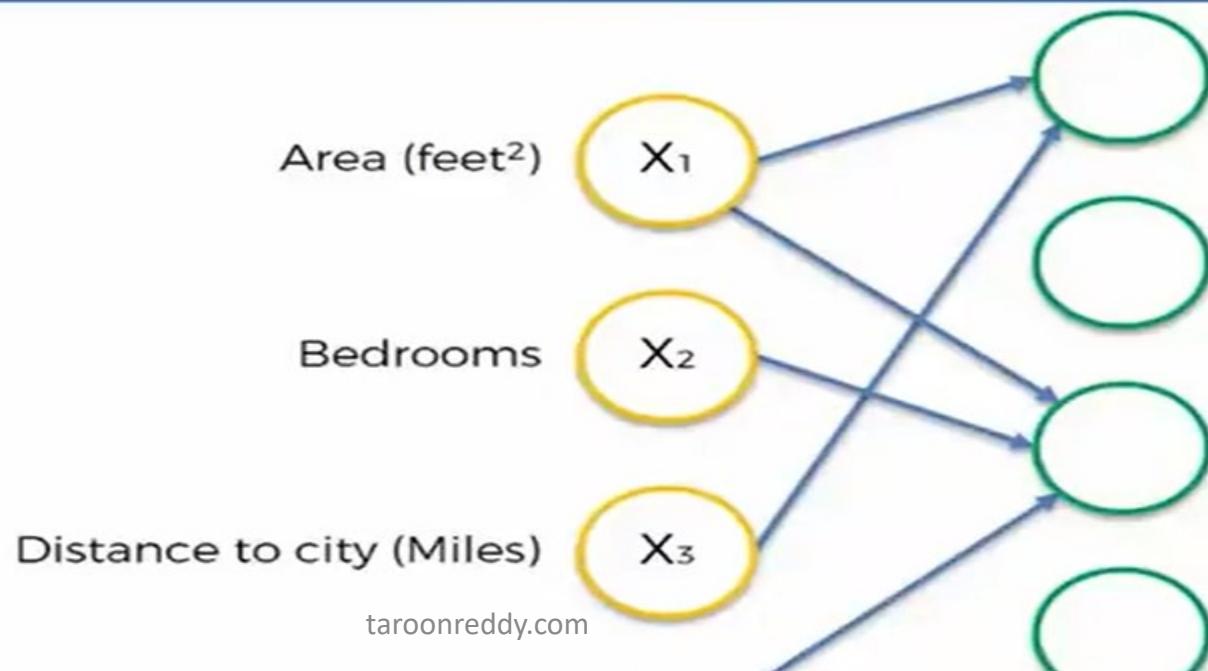
How Do Neural Networks Work?



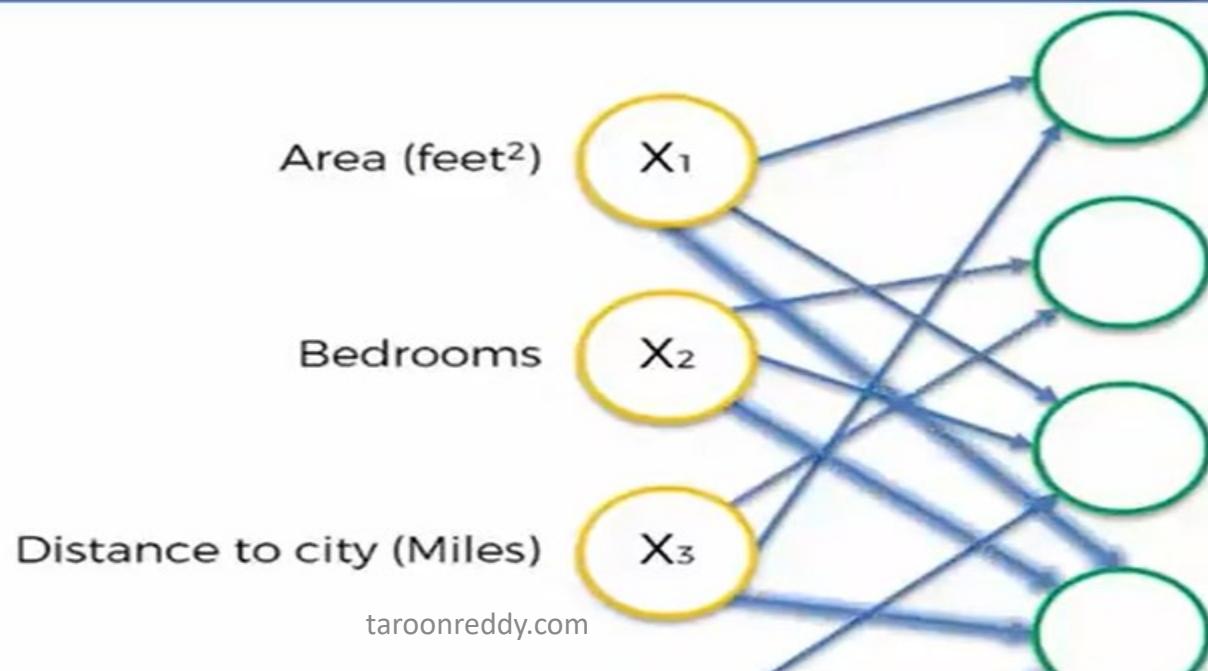
How Do Neural Networks Work?



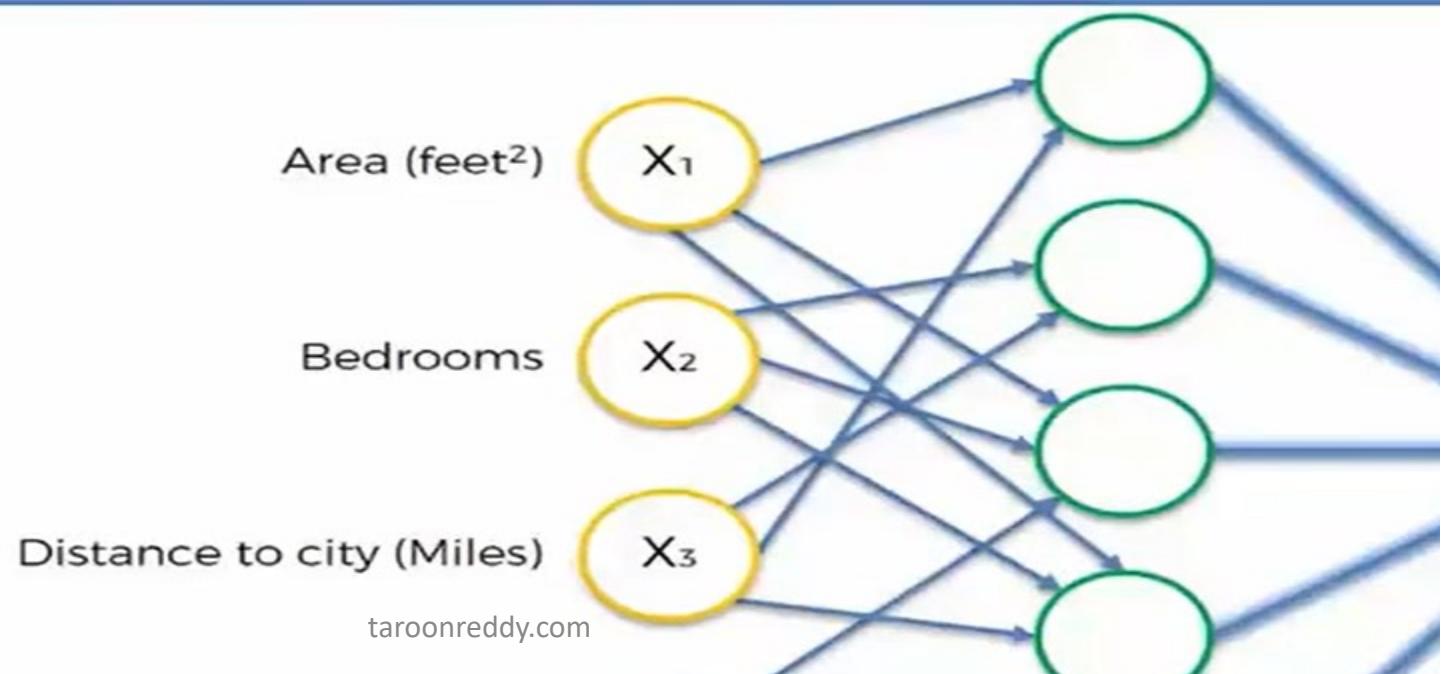
How Do Neural Networks Work?



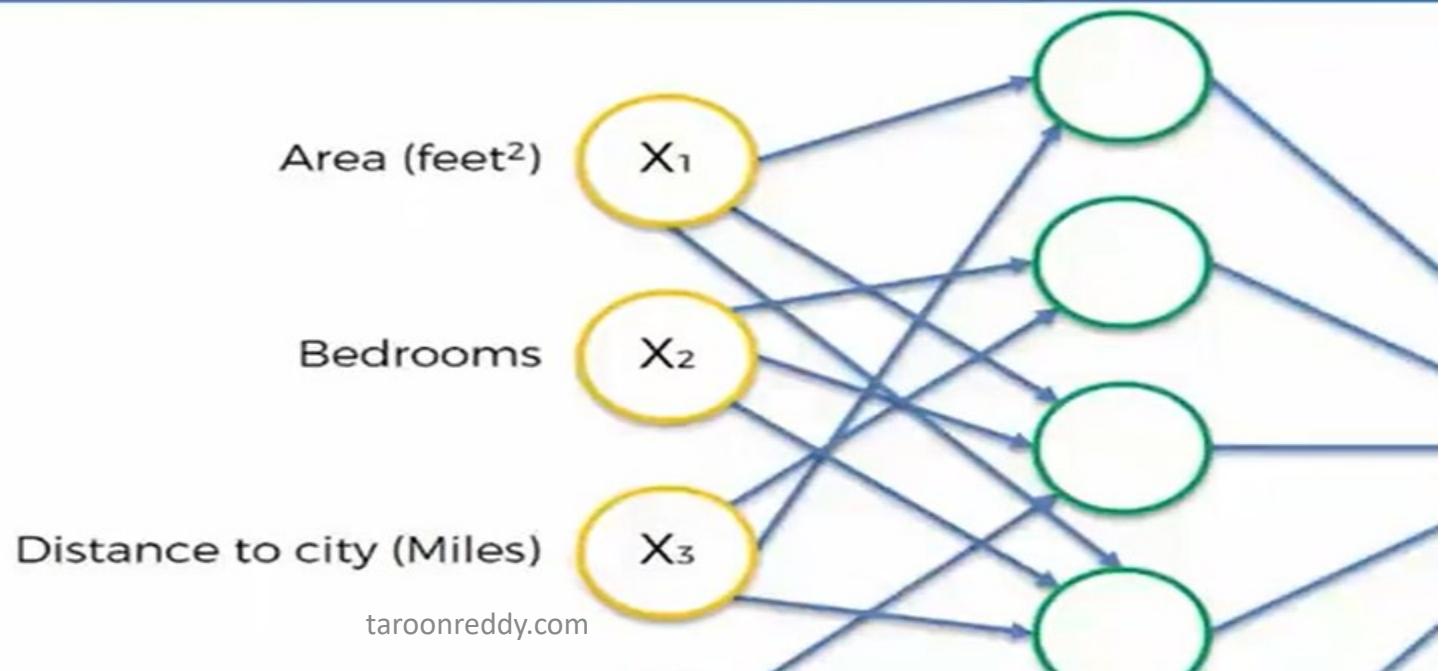
How Do Neural Networks Work?



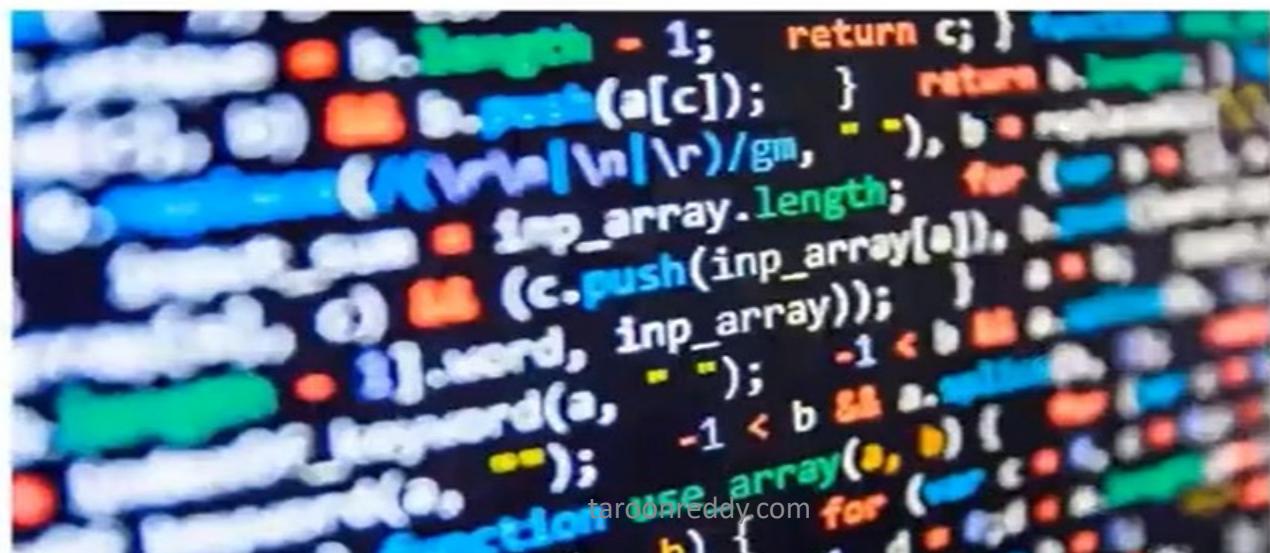
How Do Neural Networks Work?



How Do Neural Networks Work?



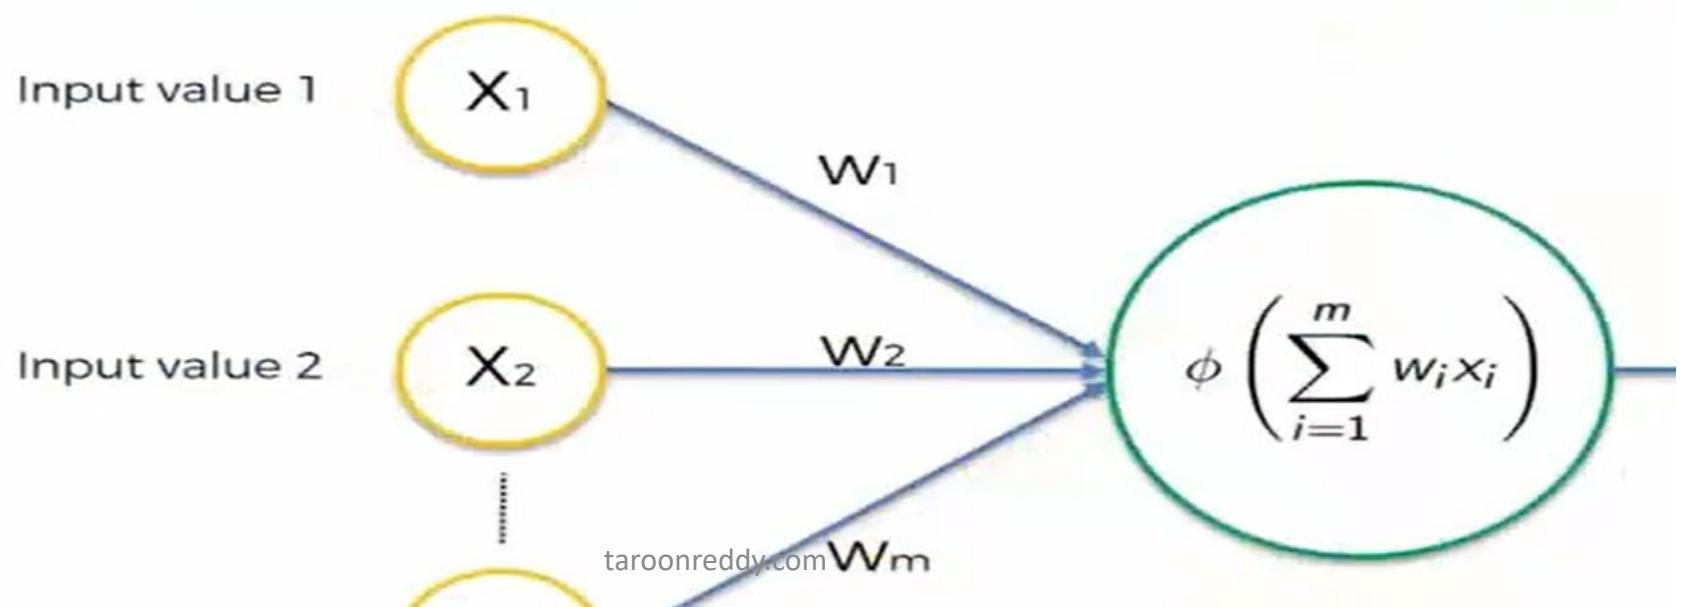
How do Neural Networ



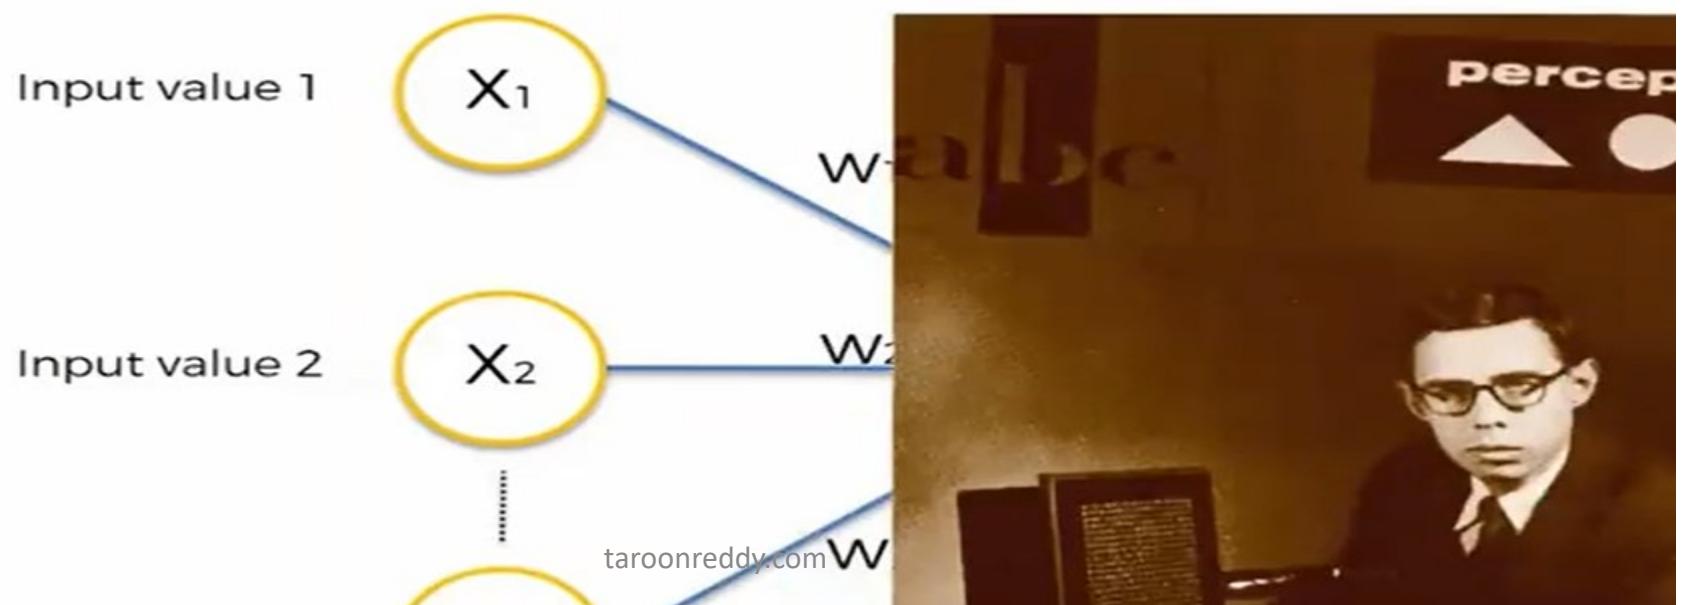
A blurred screenshot of a computer screen showing a terminal window with code and a command-line interface. The code appears to be in C or C++ and involves file operations, character processing, and array manipulations. A watermark 'taroonreddy.com' is visible at the bottom center of the image.



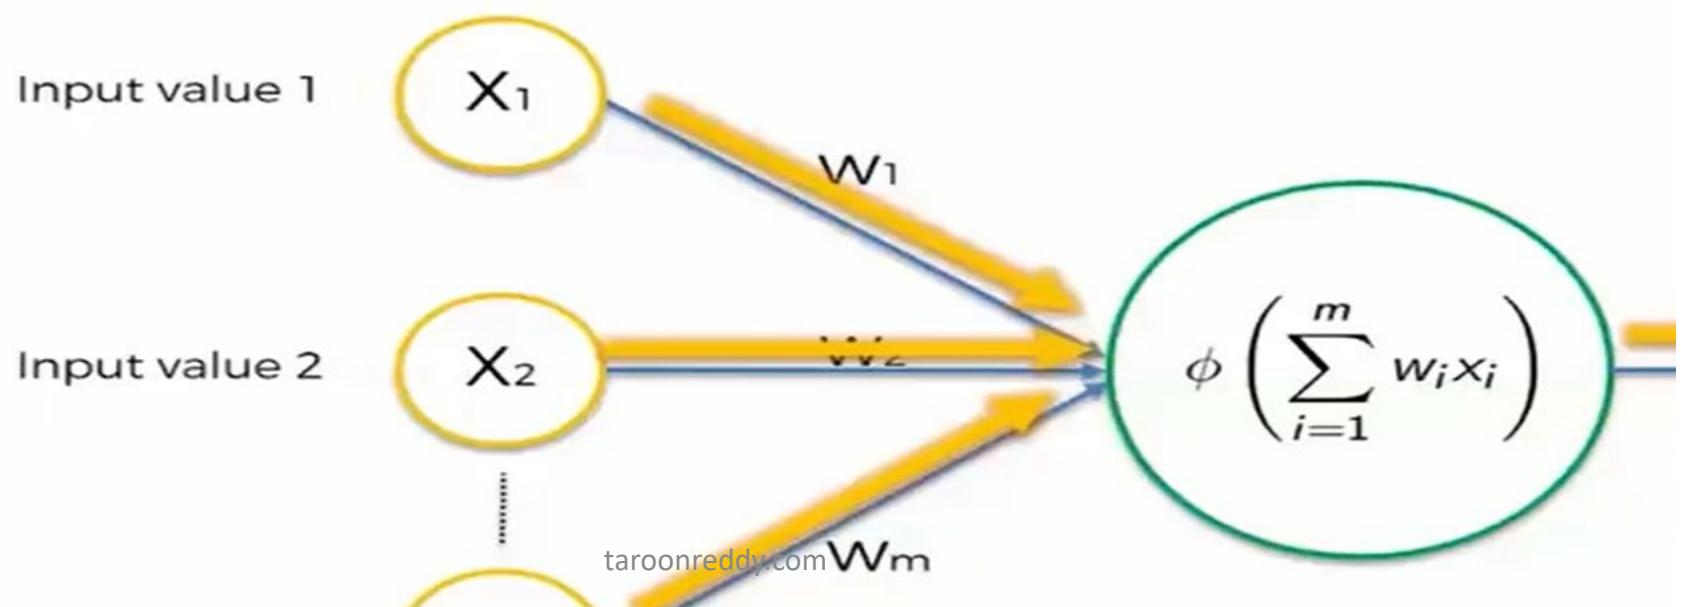
How do Neural Networks work?



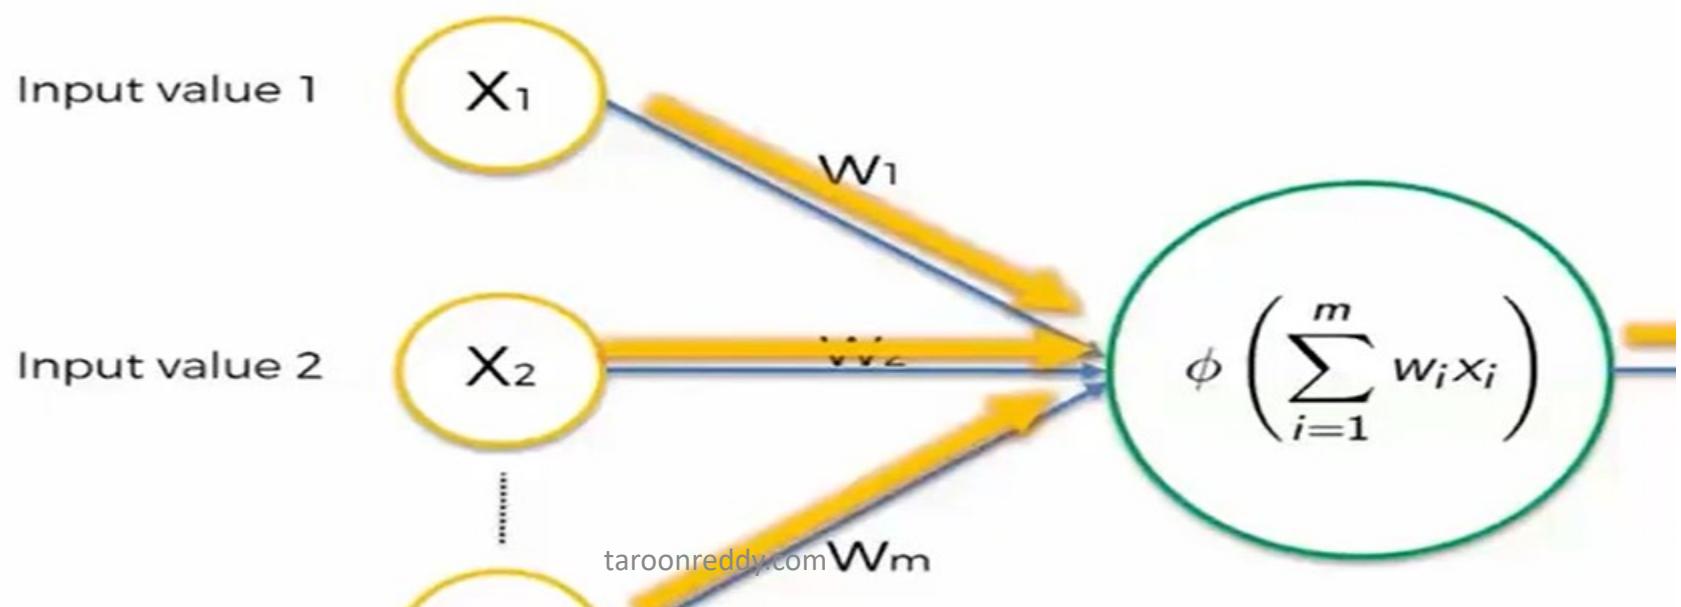
How do Neural Networks work?



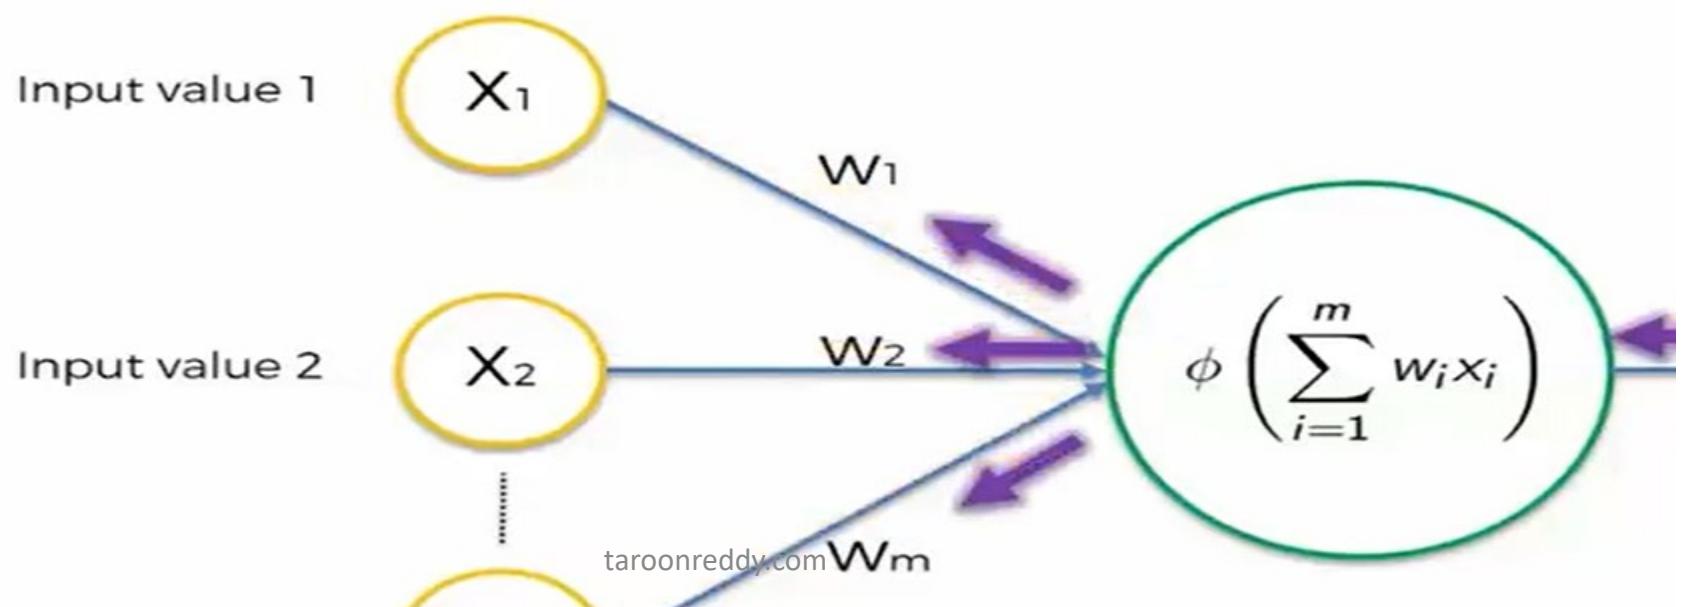
How do Neural Networks Work?



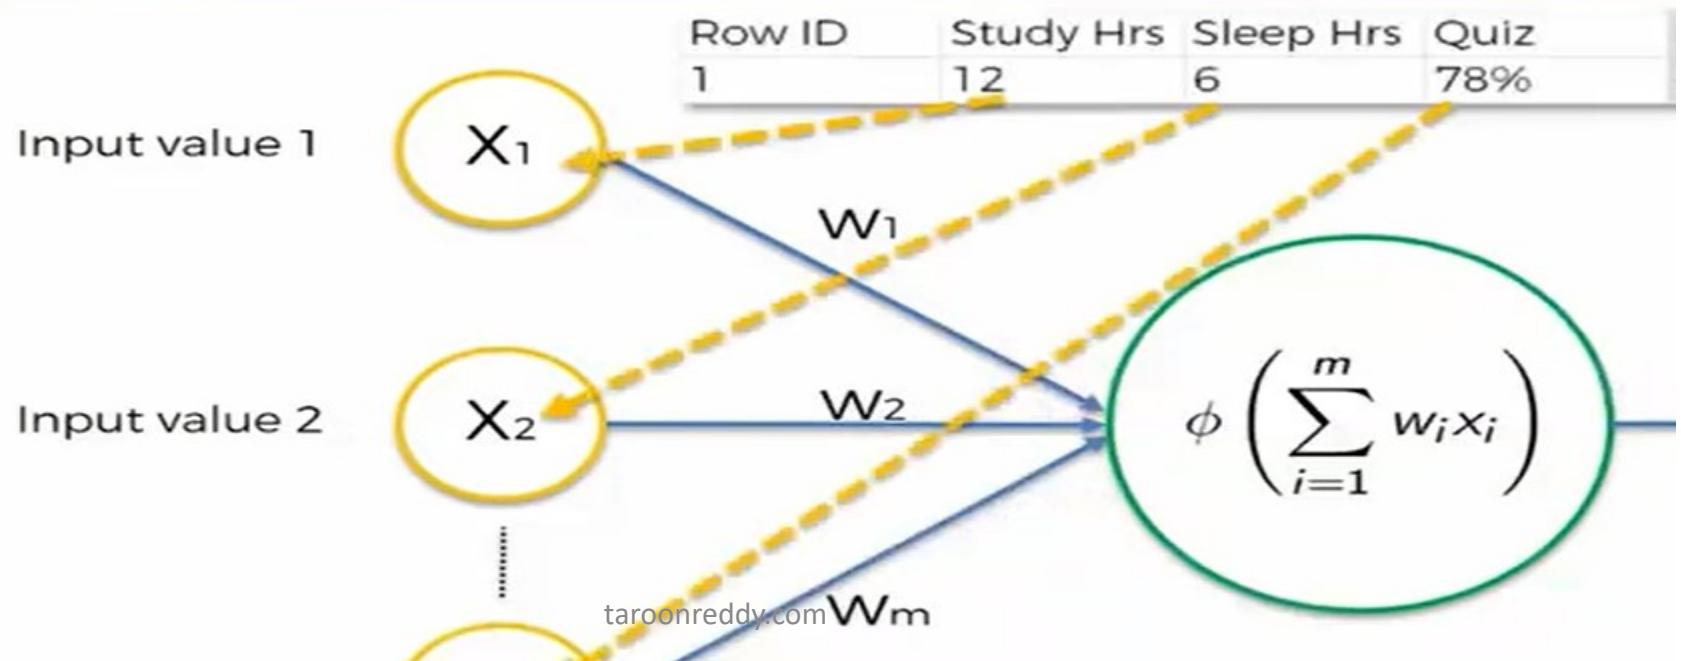
How do Neural Networks work?



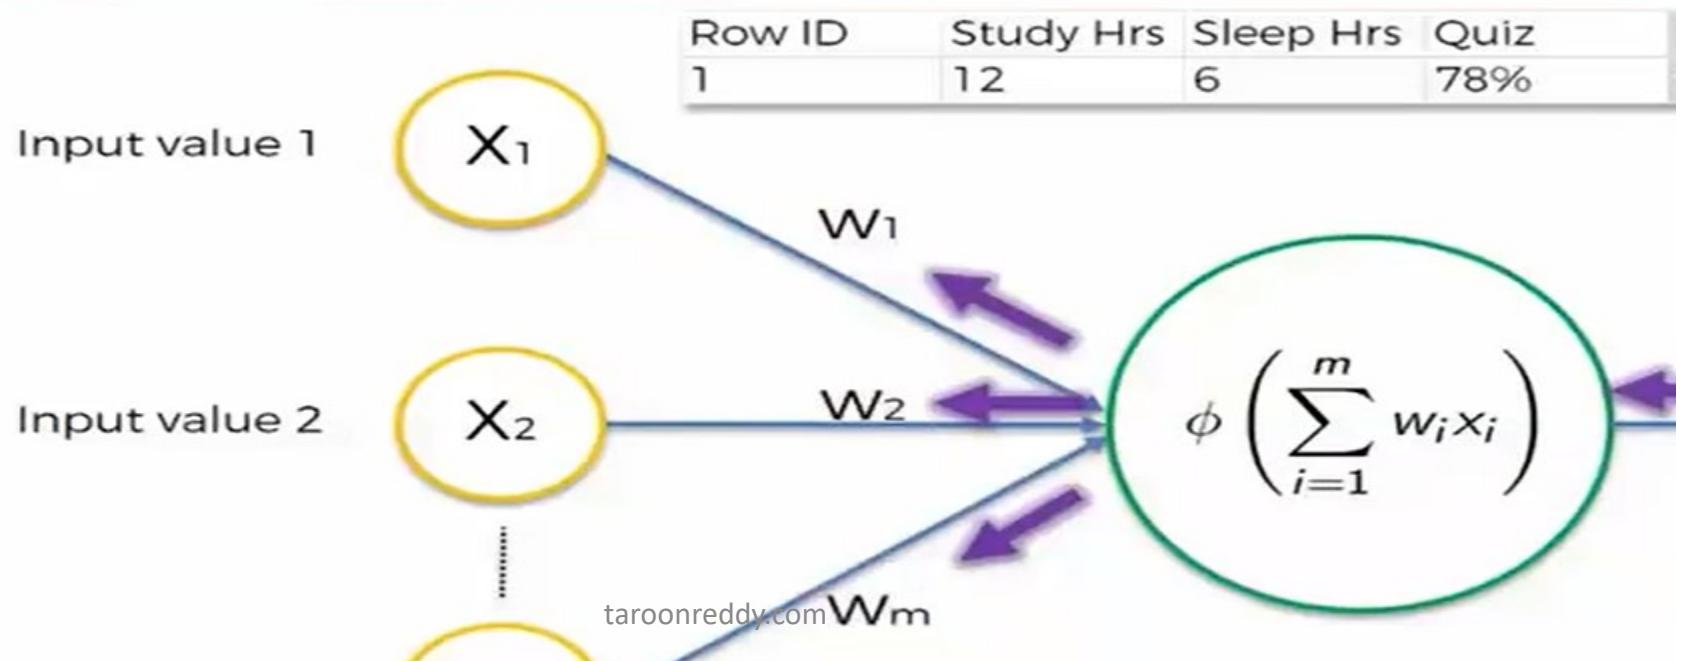
How do Neural Networks work?



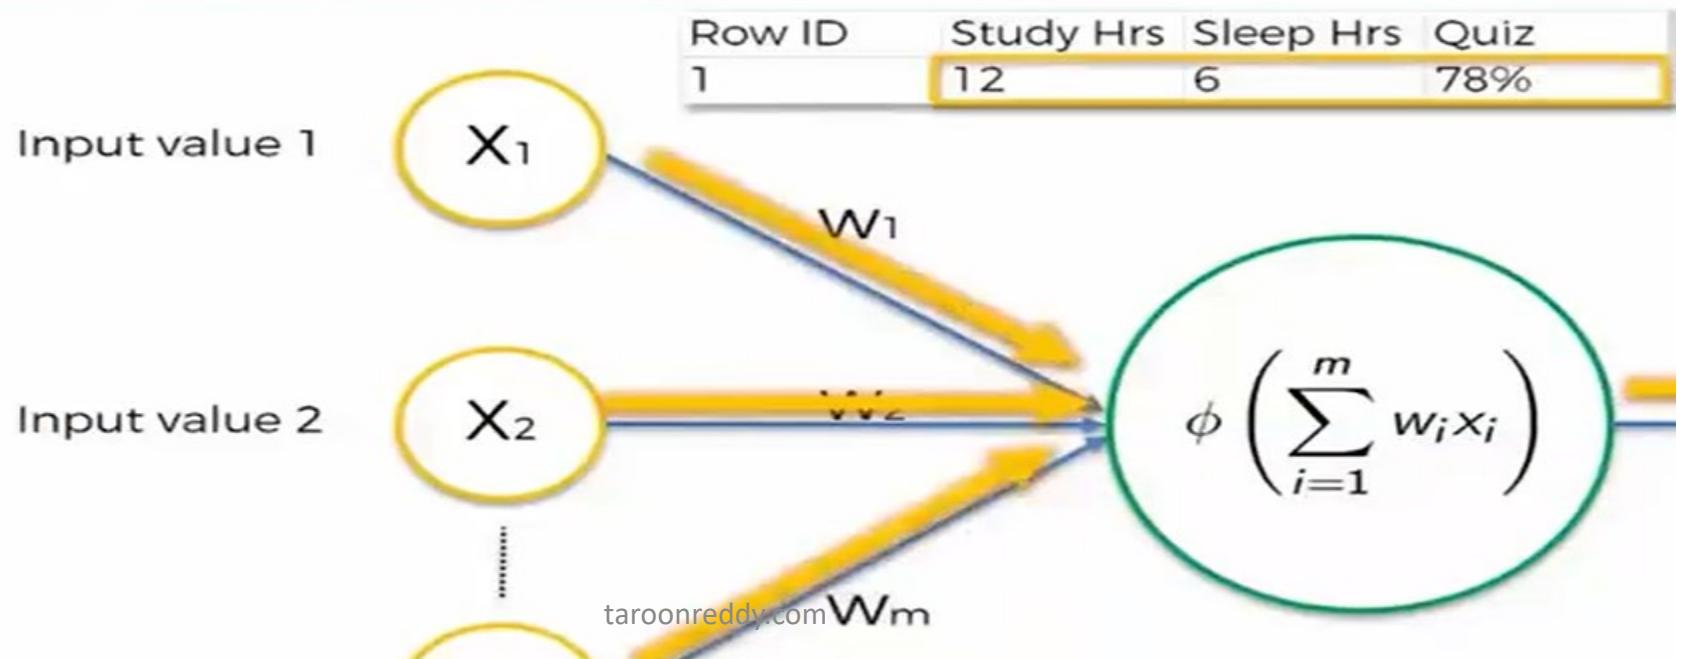
How do Neural Networks work?



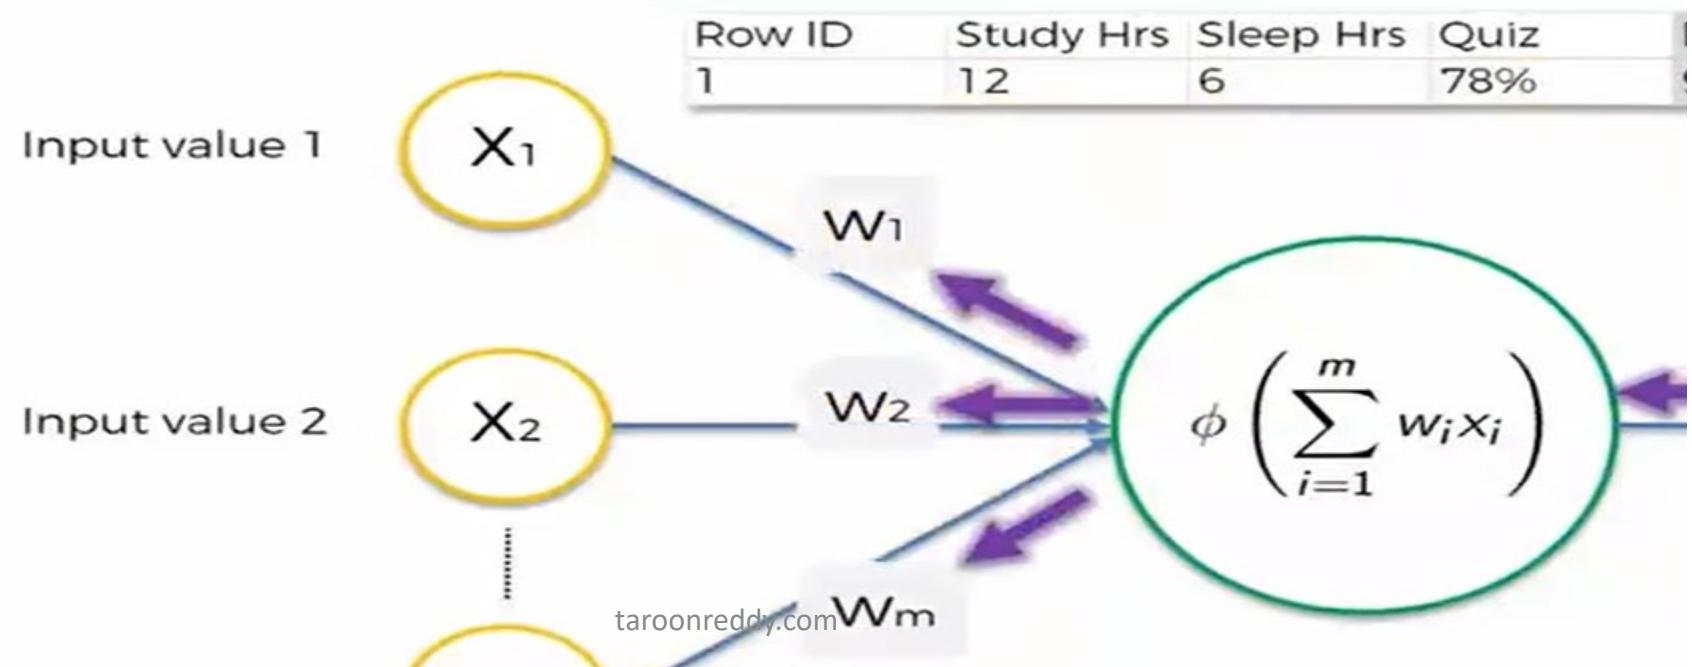
How do Neural Networks work?



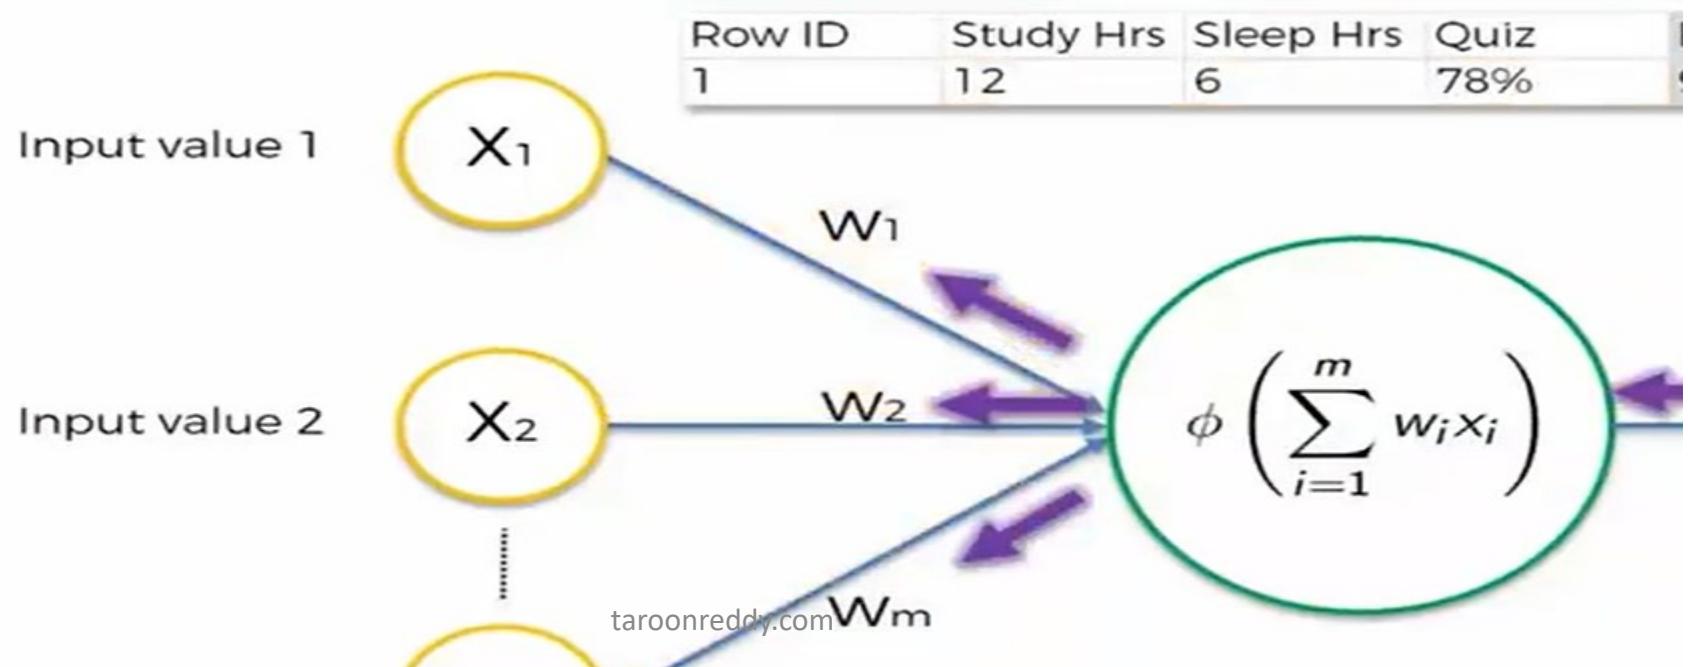
How do Neural Networks work?



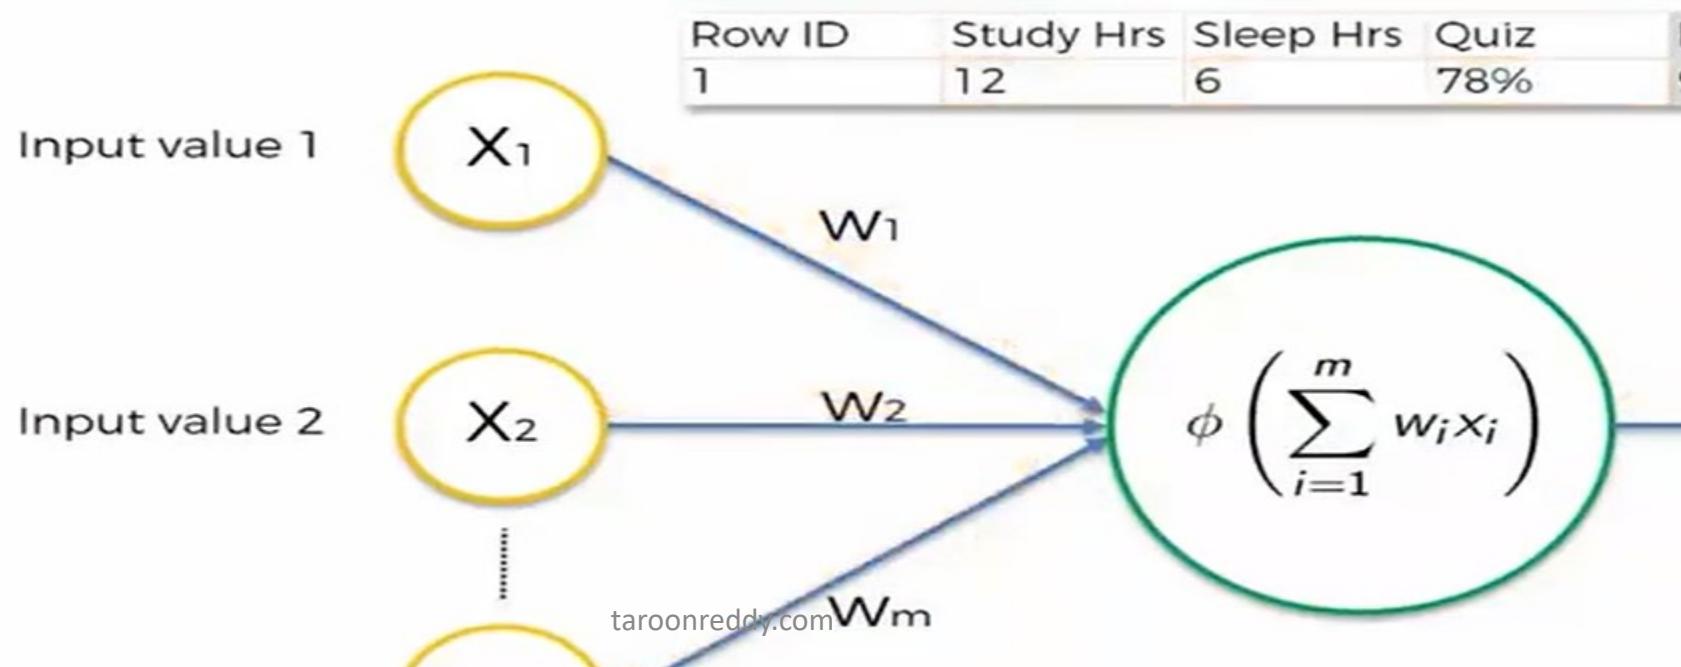
How do Neural Networks work?



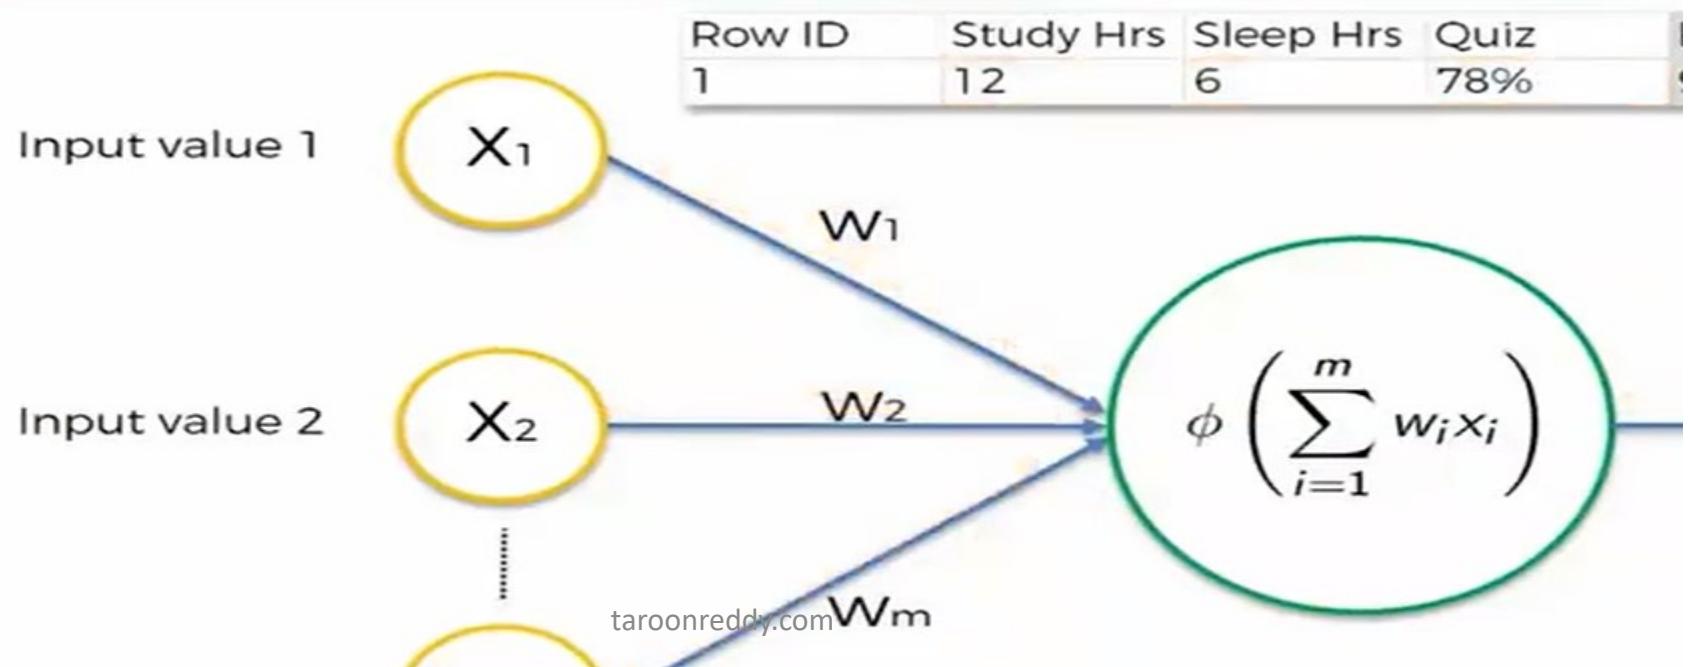
How do Neural Networks work?



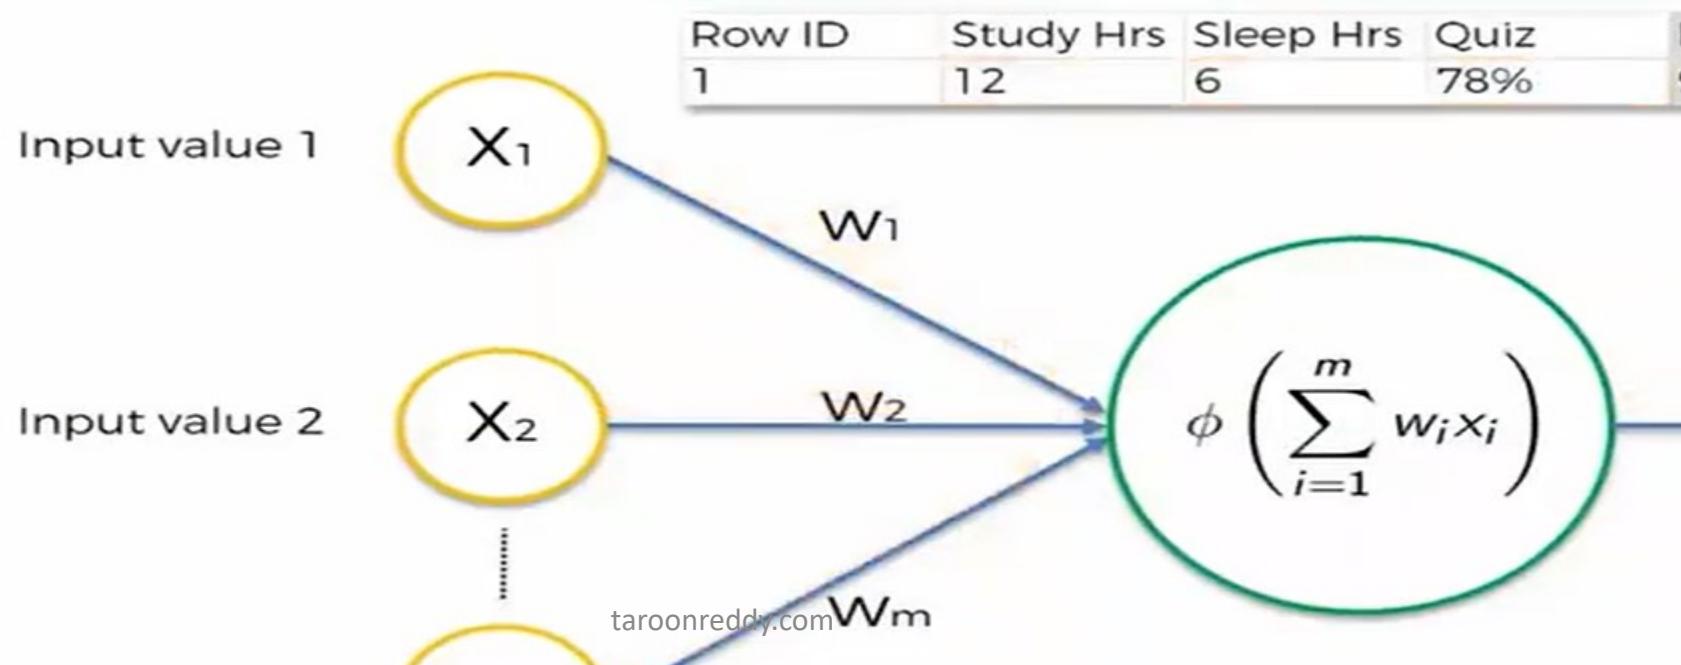
How do Neural Networks work?



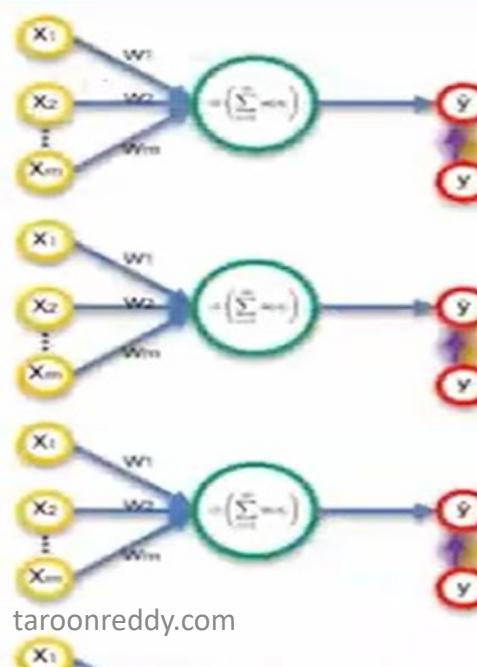
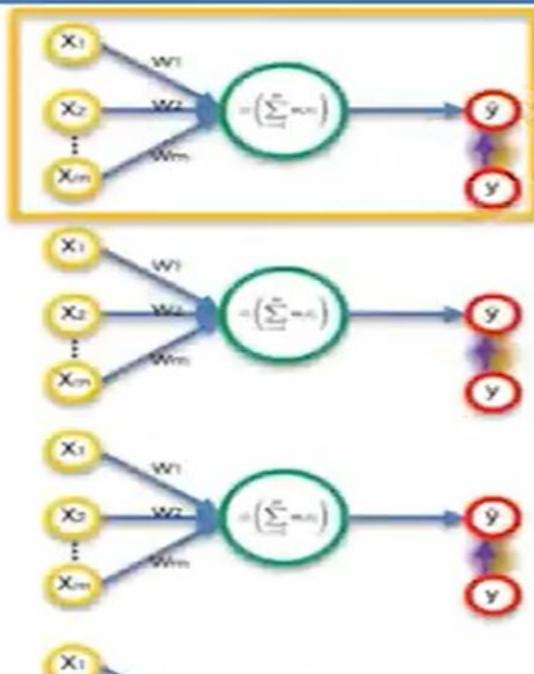
How do Neural Networks work?



How do Neural Networks work?



How do Neural Networks Work?

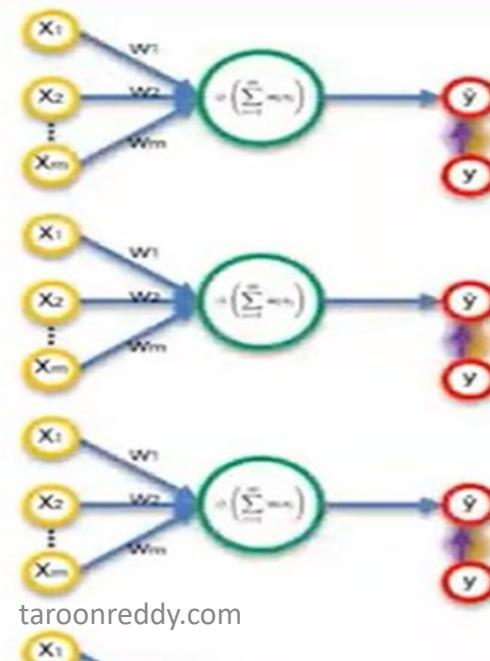
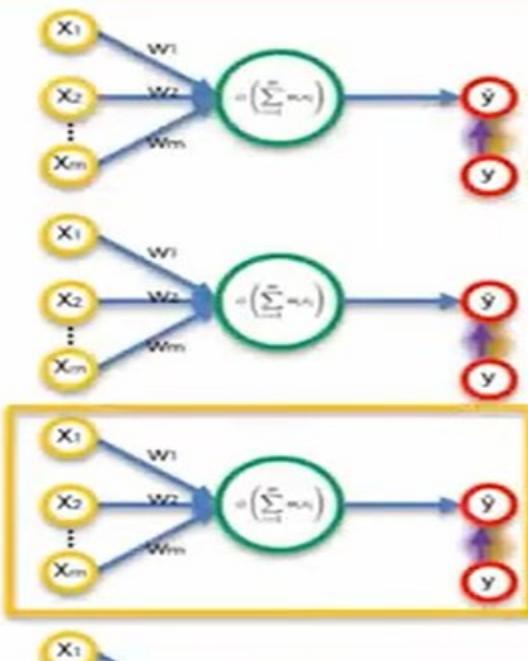


taroonreddy.com

Row ID	Study Hrs	Sle
1	12	6
2	22	6.5
3	115	4
4	31	9
5	0	10
6	5	8
7	92	6
8	57	8



How do Neural Networks Work?

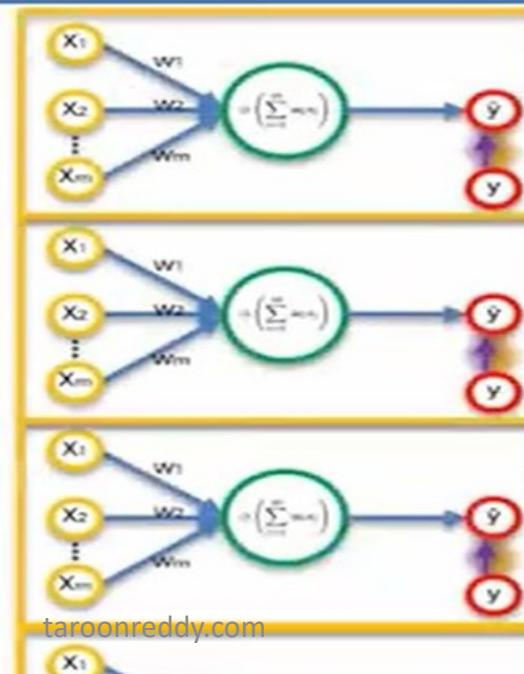
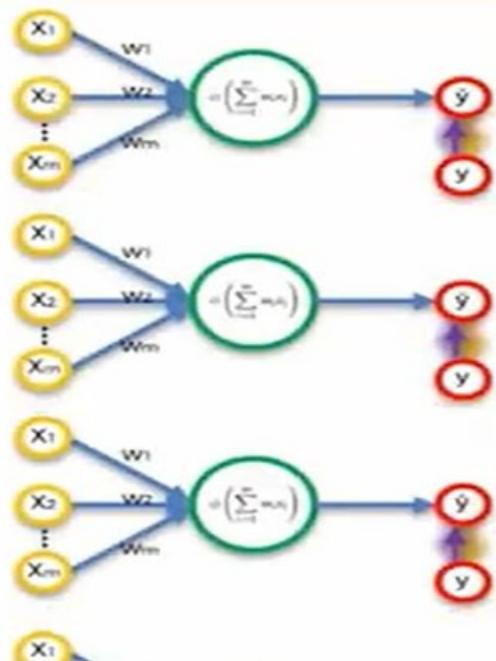


taroonreddy.com

Row ID	Study Hrs	Sle
1	12	6
2	22	6.5
3	115	4
4	31	9
5	0	10
6	5	8
7	92	6
8	57	8



How do Neural Networks Work?

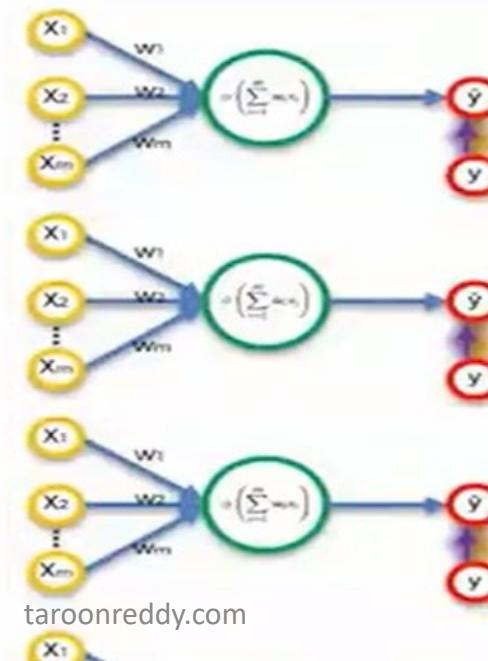
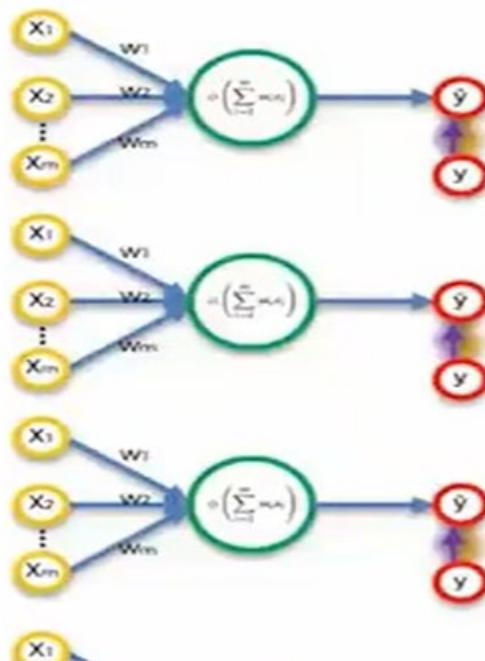


taroonreddy.com

Row ID	Study Hrs	Sle
1	12	6
2	22	6.5
3	115	4
4	31	9
5	0	10
6	5	8
7	92	6
8	57	8



How do Neural Networks Work?

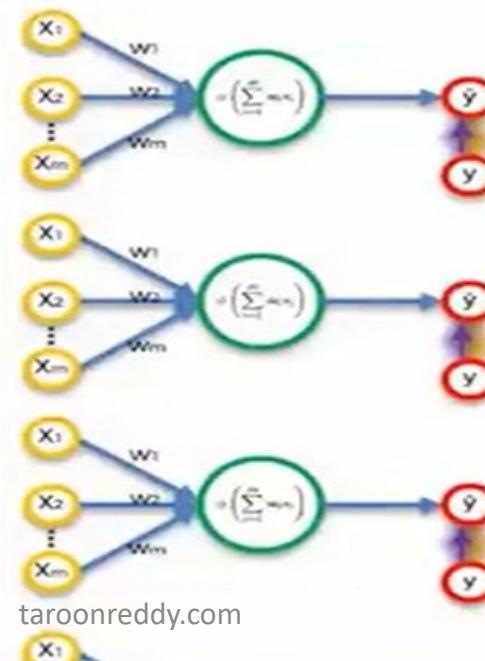
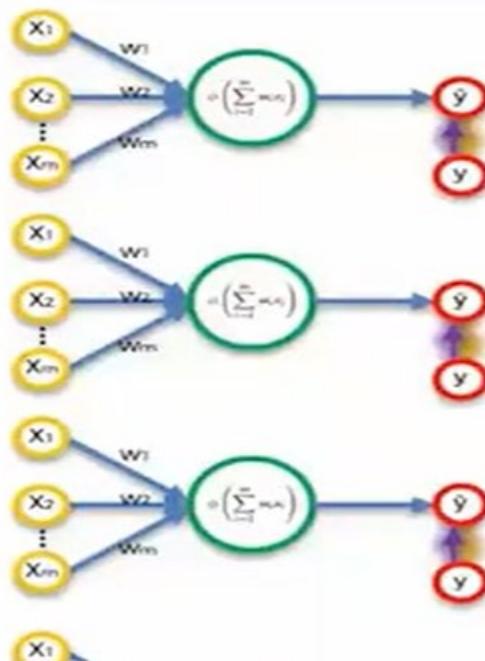


taroonreddy.com

Row ID	Study Hrs	Sle
1	12	6
2	22	6.5
3	115	4
4	31	9
5	0	10
6	5	8
7	92	6
8	57	8



How do Neural Networks Work?

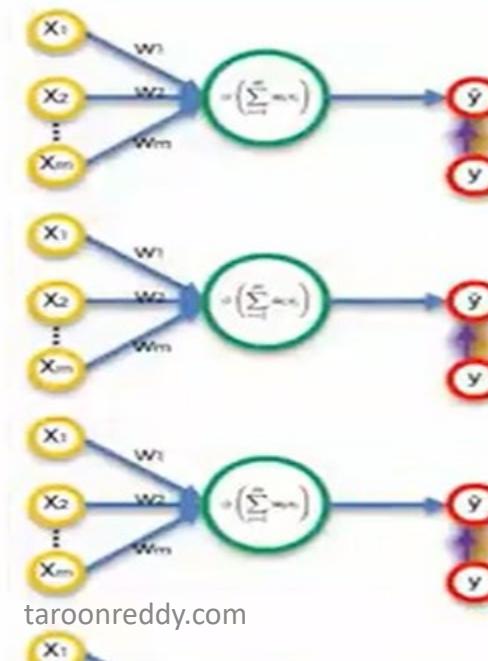
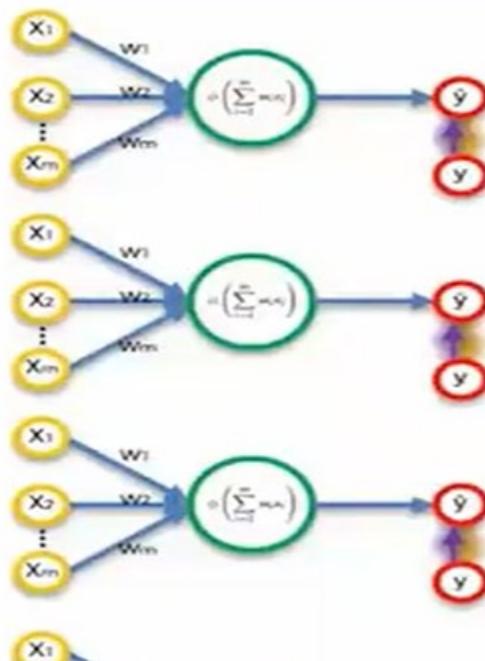


taroonreddy.com

Row ID	Study Hrs	Sle
1	12	6
2	22	6.5
3	115	4
4	31	9
5	0	10
6	5	8
7	92	6
8	57	8



How do Neural Networks Work?



Row ID	Study Hrs	Sle
1	12	6
2	22	6.5
3	115	4
4	31	9
5	0	10
6	5	8
7	92	6
8	57	8



Training the ANN with Stochastic Gradient Descent

STEP 1: Randomly initialise the weights to small numbers close to 0.



STEP 2: Input the first observation of your dataset in the input layer.



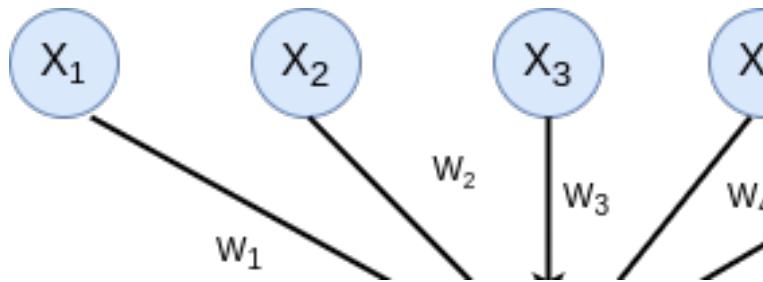
STEP 3: Forward-Propagation: from left to right, the neurons are activated. The neuron's activation is limited by the weights. Propagate the activation values to the next layer.



STEP 4: Compare the predicted result to the actual result. Measure the error.



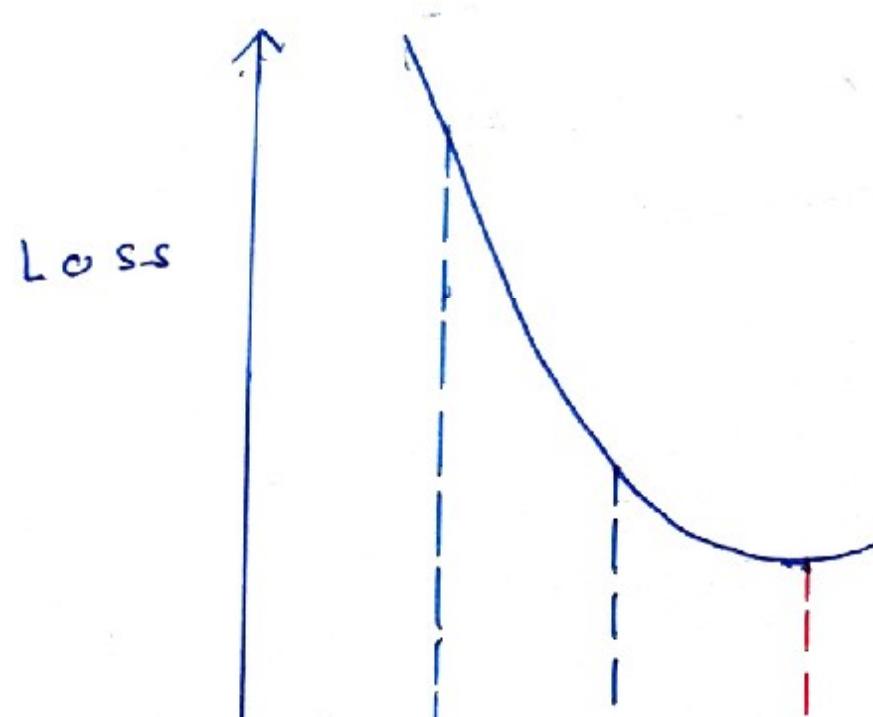
STEP 5: Back-Propagation: from right to left, the error is back-propagated through the network to determine how much they are responsible for the error. The learning rate decides how much the weights are updated.



Here w_1, w_2, w_3 are the weights of the corresponding features like x_1, x_2, x_3 and b is a constant called the bias. Its importance is that it gives flexibility. So, using such an equation the machine tries to predict a value y which may be a value we need like the price of the house. Now, the machine tries to perfect its prediction by tweaking these weights. It does so, by comparing the predicted value y with the actual value of the example in our training set and using a function of their differences. This function is called a loss function. The machine tries to decrease this loss function or the error, i.e tries to get the prediction value close to the actual value.

Gradient Descent

This method is the key to minimizing the loss function and achieving our target, which is to predict close to the original value.



In this diagram, above we see our loss function graph. If we observe we will see it is basically a parabolic shape or a convex shape, it has a specific global minimum which we need to find in order to find the minimum loss function value. So, we always try to use a loss function which is convex in shape in order to get a proper minimum. Now, we see the predicted results depend on the weights from the equation. If we replace equation 1 in equation 2 we obtain this graph, with weights in X-axis and Loss on Y-axis.

Initially, the model assigns random weights to the features. So, say it initializes the weight= a . So, we can see it generates a loss which is far from the minimum point L_{min} . Now, we can see that if we move the weights more towards the positive x-axis we can optimize the loss function and achieve minimum value. But, how will the machine know? We need to optimize weight to minimize error, so, obviously, we need to check how the error varies with the weights. To do this we need to find the derivative of the Error with respect to the weight. **This derivative is called Gradient.**

$$\text{Gradient} = dE/dw$$

Where E is the error and w is the weight.

Let's see how this works. Say, **if the loss increases with an increase in weight so Gradient will be positive**, So we are basically at the point C, where we can see this statement is true. **If loss decreases with an increase in weight so gradient will be negative**. We can see point A, corresponds to such a situation. Now, from point A we need to move towards positive x-axis and the gradient is negative. From point C, we need to move towards negative x-axis but the gradient is positive. **So, always the negative of the Gradient shows the directions along which the weights should be moved in order to optimize the loss function**. So, this way the gradient guides the model whether to increase or decrease weights in order to optimize the loss function.

The model found which way to move, now the model needs to find by how much it should move the weights. This is decided by a parameter called **Learning Rate denoted by Alpha**. In the diagram we see, the weights are moved from point A to point B which are at a distance of dx .

$$dx = \alpha * |dE/dw|$$

So, the distance to move is the product of learning rate parameter alpha and the magnitude of change in error with a change in weight at that point.

Now, we need to decide the Learning Rate very carefully. If it is very large the values of weights will be changed with a great amount and it would overstep the optimal value. If it is very low it takes tiny steps and takes a lot of steps to optimize. The updated weights are changed according to the following formula.

$$w = w - \alpha * |dE/dw|$$

where w is the previous weight.

With each epoch, the model moves the weights according to the gradient to find the best weights.

Now, this is a loss optimization for a particular example in our training dataset. Our dataset contains thousands of such examples, so it will take a huge time to find optimal weights for all. Experiments have shown that if we optimize on only one sample of our training set, the weight optimization is good enough for the whole dataset. So, depending upon the methods we have different types of gradient descent mechanisms.

Gradient Descent Methods

Stochastic Gradient Descent: When we train the model to optimize the loss function using only one particular example from our dataset, it is called Stochastic Gradient Descent.

Batch Gradient Descent: When we train the model to optimize the loss function using the mean of all the individual losses in our whole dataset, it is called Batch Gradient Descent.

Mini-Batch Gradient Descent: Now, as we discussed batch gradient descent takes a lot of time and is therefore somewhat inefficient. If we look at SGD, it is trained using only 1 example. So, how good do you think a baby will learn if it is shown only one bike and told to learn about all other bikes? It's simple its decision will be somewhat biased to the peculiarities of the shown example. So, it is the same for the SGD, there is a possibility that the model may get too biased with the peculiarity of that particular example. So, we use the mean of a batch of 10–1000 examples to check the optimize the loss in order to deal with the problems.

Keras & Tensorflow

taroonreddy.com

Artificial Neural Network

taroonreddy.com